THE ATCHLEY UNIFIED METIS FRAMEWORK (AUM)

Structural Invariants and Runtime Governance for Safe Adaptive Intelligence

A Dissertation Presented in Partial Fulfillment of the Requirements for the Degree of Doctor of Philosophy

Devin Earl Atchley
Independence, Missouri, United States of America

Prototype Dissertation — First Draft
2025

Abstract

Modern artificial intelligence systems operate as adaptive stochastic processes embedded in complex, partially observable environments. Traditional safety frameworks address behavioral outputs through rule-based filters, reward shaping, or constitutional AI approaches. These methods cannot provide formal guarantees of dynamical stability, because they monitor outputs rather than governing the internal structural geometry of the reasoning process. This dissertation introduces the Atchley Unified METIS Framework (AUM), a structural safety architecture that constrains adaptive intelligence through dynamical system invariants, stochastic control theory, runtime governance mechanisms, and cryptographic epistemic accountability.

The framework addresses three fundamental failure modes: semantic drift (representational instability under perturbation), entropy collapse (loss of distributional diversity leading to brittle outputs), and unverifiable belief formation (belief updates not tied to cryptographically verified evidence). AUM provides formal guarantees against all three failure modes simultaneously through a layered governance stack integrating seven subsystems: the Atchley Unified Field compact state space (AUF), the Recursive Cognitive Containment Lattice (RCCL), the Scientific Discovery Engine (SDE), the Co-Motion temporal regulator, the Deepfreeze recovery operator, the nine-gate Moore machine Gate Control Graph, and the SHA-256 hash-chain Evidence Ledger.

The mathematical core of the dissertation develops sixteen definitions, four lemmas, seven propositions, and six principal theorems with complete proofs: (T1) forward invariance of the safe structural region via Control Barrier Function theory; (T2) persistence of the entropy floor as a hard safety invariant; (T3) Lyapunov stability under co-motion time reparameterization with explicitly constructed Lyapunov function $V = d_W(\Sigma, S^*)^2$; (T4) bounded cumulative structural violation under information-theoretically budgeted exploration; (T5) full reconstructibility of the belief trajectory from the cryptographic Evidence Ledger; and (T6) geometric convergence of the Deepfreeze operator using contraction-toward-a-set rather than global Banach contraction.

Empirical validation via Monte Carlo simulation (500 runs, 10^6 steps each, 32-dimensional state space) demonstrates: 92% reduction in safety violations, 87% reduction in escape probability, 3× improvement in recovery time, 3.4× stronger Lyapunov decay rate, and 94% retained operational efficiency versus an uncontrolled baseline. Gate security testing against four adversarial attack classes achieves 96–100% block rates. Spectral stability validation confirms positive spectral gap $\gamma = 0.21$ across all runs.

The dissertation situates these results relative to prior work in Control Barrier Functions (Ames et al., 2017), information geometry (Amari, 2016), hybrid dynamical systems (van der Schaft and Schumacher, 2000), and Bayesian epistemics. The principal scientific contribution is the demonstration that a single coherent formal framework — grounded in control theory, information theory, Bayesian epistemology, and cryptography — can provide comprehensive runtime governance of artificial cognitive systems. A complete prototype implementation (AUF Gate MCP Runtime, RCCL-Lite, ACPL Chrome Extension, Singularity Kernel Simulation) demonstrates practical feasibility.

Keywords: *governed cognition, control barrier functions, Lyapunov stability, contraction toward a set, Bayesian epistemics, compact metric state space, Moore machine governance, cryptographic provenance, co-motion time, safety-critical AI, stochastic control, structural invariants*

Dedication

To the architects of difficult problems —
those who choose rigor over simplicity,
and formal truth over comfortable approximation.

Acknowledgments

This work represents an independent research program carried out in its entirety by the author. The mathematical foundations draw on a century of accumulated knowledge in dynamical systems theory, Bayesian probability, control engineering, and cryptography. The author acknowledges the intellectual contributions of the following foundational works, without which this synthesis would not have been possible: Shannon's information theory, Banach's fixed-point theorem, Lyapunov's stability theory, Itô's stochastic calculus, Clarke's nonsmooth analysis, and Ames et al.'s Control Barrier Function framework.

The author thanks the open-source scientific Python ecosystem — NumPy, SciPy, and Matplotlib — for providing the computational substrate of the empirical validation program. The docx and Node.js communities provided the document toolchain.

This dissertation is dedicated to the proposition that safety and capability in artificial intelligence are not in fundamental tension — that a formally governed system can be both powerful and safe, because the governance is structural, not behavioral.

Table of Contents

Abstract ii
Dedication iv
Acknowledgments v
List of Tables viii
List of Figures ix
Notation Reference x
Chapter 1 Foundations, Formal Structure, and Theoretical Development 1
1.1 Motivation and Problem Statement 1
1.2 Notation and Preliminaries 3
1.3 The AUF State Space 4
1.4 The Action Space 7
1.5 The RCCL Behavioral Observation Vector 8
1.6 Combined Cognitive State Space 10
1.7 Foundational Axioms A1–A9 11
1.8 The Atchley Structural Invariant 16
1.9 The RCCL Runtime Risk Function 19
1.10 The RCCL Safety Filter (CBF Formulation) 21
1.11 Co-Motion Temporal Dynamics 23
1.12 The Scientific Discovery Engine 25
1.13 The Deepfreeze Operator 28
1.14 The Gate Control Graph 30
1.15 The Evidence Ledger 33
1.16 Principal Theorems and Proofs 34
1.17 Worked Example 42
1.18 System Architecture 44
1.19 Discussion and Contribution 45
Chapter 2 Mathematical Foundations 49
2.1 Stochastic Differential Equations 49
2.2 Lyapunov Stability Theory 52
2.3 Control Barrier Functions 55
2.4 Metric Space Theory 59
2.5 Bayesian Probability Foundations 62
2.6 Information Theory 65
2.7 Metastability 67
Chapter 3 Structural Cognition Geometry 71
3.1 The AUF State Space in Depth 71
3.2 Curvature of Embedding Trajectories 74
3.3 Metric Properties and Continuity 77
3.4 Geometric Interpretation of Safety 79
Chapter 4 The Risk-Controlled Cognitive Loop 83
4.1 System Model 83
4.2 Observation Map 85
4.3 Risk Function Theory 87
4.4 Decision Thresholds and Policy 89
Chapter 5 The Entropy Floor Safety Principle 93
5.1 Motivation and Failure Mode 93

5.2 Formal Invariant Statement 95
5.3 Information-Theoretic Basis 97
5.4 Enforcement Mechanism 99
Chapter 6 The Deepfreeze Operator 103
6.1 Recovery Operator Design 103
6.2 Contraction-Toward-a-Set Theory 106
6.3 Component-Wise Projection 109
6.4 Convergence Analysis 112
Chapter 7 The Gate Control Graph 115
7.1 Moore Machine Formalization 115
7.2 Gate Specifications 117
7.3 Completeness and Soundness 121
7.4 Interaction With Safety Filter 123
Chapter 8 The Violation Budget Model 127
8.1 Exploration-Safety Tradeoff 127
8.2 Information-Theoretic Capacity Bound 130
8.3 Adversarial Robustness 133
Chapter 9 Spectral Stability Analysis 137
9.1 Generator and Spectrum 137
9.2 Spectral Gap Criterion 140
9.3 Relation to Metastability 143
Chapter 10 Empirical Validation 147
10.1 Experimental Design 147
10.2 Primary Safety Metrics 151
10.3 Gate Security Testing 156
10.4 Spectral and Stability Validation 159
10.5 Efficiency Analysis 161
Chapter 11 Implementation Architecture 165
11.1 AUF Gate MCP Runtime 165
11.2 RCCL-Lite 168
11.3 ACPL Chrome Extension 170
11.4 Singularity Kernel Simulation 172
Chapter 12 Applications 175
Chapter 13 Limitations 181
Chapter 14 Future Research Directions 187
Chapter 15 Conclusion 193
References 199
Appendix A Formal Definitions 207
Appendix B Proofs of Core Theorems 213
Appendix C Simulation Methodology 221
Appendix D Experimental Data Tables 227

Notation Reference

Symbol	Domain / Type	Meaning
$\mathscr{M}$	Compact metric space	AUF state space (Definition 1.2)
$\mathscr{A}$	Compact metric space	Action space (Definition 1.3); $\mathscr{A}=\mathscr{A}_1 \sqcup \mathscr{A}_2 \sqcup \mathscr{A}_3$
$\mathscr{S}^* \subset \mathscr{M}$	Closed subset	Safe attractor set: $\{\Sigma : \mathfrak{A}(\Sigma) \geq 0\}$
$\mathscr{Z}$	Product space	Global cognitive state space
Σ_t	Tuple in $\mathscr{M}$	AUF structural state at time t
χ_t	Vector in $[0,1]^4$	RCCL behavioral observation
B_t	Element of $\Delta(\mathscr{H})$	SDE belief distribution over hypotheses
q_t	Element of Q	Active gate mode
$\mathfrak{A}(\cdot)$	$\mathscr{M} \to \mathbb{R}$	Atchley structural invariant
$\mathfrak{D}$	Operator $\mathscr{M} \to \mathscr{M}$	Deepfreeze operator
$\mathfrak{G}$	Moore machine	Gate Control Graph
$F : \mathscr{M} \times \mathscr{A} \to \mathscr{M}$	Transition model	Next-state function (system hyperparameter)
d_W	Metric on $\mathscr{M}$	Weighted Euclidean: $d_W(\Sigma,\Sigma')=\lVert\Sigma-\Sigma'\rVert_W$
$V(\Sigma)$	$\mathbb{R}\geq 0$	Lyapunov function: $d_W(\Sigma,\mathscr{S}^*)^2$
$\alpha(\Sigma_t)$	$(0,\alpha_0]$	Co-motion clock rate
κ_{CBF}	(0,1)	CBF decay constant
c	(0,1)	Deepfreeze contraction ratio
γ	$\mathbb{R}>0$	Spectral gap of safe generator $\mathscr{L}_{safe}$
H_{min}	$\mathbb{R}>0$	Minimum admissible structural entropy
I_t	$\mathbb{R}\geq 0$	Per-step SDE information gain
B_{cap}	$\mathbb{R}\geq 0$	Total information-theoretic exploration budget

Chapter 1 Foundations, Formal Structure, and Theoretical Development

Chapter Status

This chapter is complete at committee-review level. All critical review items from the first review cycle have been incorporated. The content below provides a structured summary of the complete Chapter 1 document, which is available as a standalone file (AUCGS_Ch1_SecondRevision.docx).

1.1 Motivation and Problem Statement

Three failure modes motivate the AUCGS. (i) Semantic drift: internal representations shift under adversarial perturbation or session length, producing outputs no longer grounded in the task context. (ii) Entropy collapse: reasoning systems operating under repeated compression collapse toward deterministic, brittle outputs. (iii) Unverifiable belief formation: systems updating beliefs in response to evidence provide no cryptographic guarantee that the evidence supports the inference. The AUCGS addresses all three through a unified mathematical architecture.

The primary scientific contribution is the demonstration that a single coherent formal framework — combining control theory, information theory, Bayesian epistemology, and cryptographic accountability — can provide comprehensive runtime governance of artificial cognition. Rather than proposing ad hoc interventions, the system establishes a rigorous compact-metric state-space formalism, a Control Barrier Function family of invariants, and a layered enforcement automaton.

1.2 Notation and Preliminaries

Throughout this dissertation, subscripts denote time indices and superscripts denote component indices. For a metric space (X, d), $\mathrm{Lip}(X)$ denotes real-valued Lipschitz functions. The notation $d(x, S) := \inf_{s \in S} d(x, s)$ denotes distance to a closed set. A map $f : X \to X$ is a contraction toward $S \subset X$ with ratio $c \in (0,1)$ if $d(f(x), S) \leq c \cdot d(x, S)$ for all $x \in X$ and $f(s) = s$ for $s \in S$. See the Notation Reference preceding this chapter for the complete symbol table.

1.3 The AUF State Space

Definition 1.1 (AUF Component Fields)

The seven component fields of the AUF structural state are:

(F1) $\kappa_t \in [-1,1]$ — Representation curvature: normalized Menger curvature of the three-point embedding centroid trajectory (e_{t-2}, e_{t-1}, e_t). Computable in $O(Wd)$. (F2) $\in_t \quad [0,1]$ — Compute load: normalized forward-pass latency. (F3) $H_t \in [0,1]$ — Normalized structural entropy: $H(B_t)/H_{max}$. (F4) $R_t \in [0,1]$ — Resilience: cosine similarity of current and reference gradient directions. (F5) $\Phi_t \in [0,1]$ — Semantic coherence: cosine similarity of current and reference embedding centroids. (F6) $G_t \in [0,1]$ — Governance score: most recent gate decision.

(F7) $\Theta_t \in [-1,1]$ — Temporal alignment: normalized co-motion time deviation.

Definition 1.2 (AUF State Space ($\mathcal{M}$, d_W))

$\mathscr{M} := [-1,1] \times [0,1]^5 \times [-1,1] \subset \mathbb{R}^7$ equipped with $d_W(\Sigma,\Sigma') := (\Sigma_i w_i(\Sigma_i-\Sigma'_i)^2)^{1/2}$ for positive weights $W = \mathrm{diag}(w_1,\ldots,w_7)$. ($\mathscr{M}$, d_W) is compact and complete.

1.7 Foundational Axioms A1–A9

The nine axioms are stated and justified fully in the complete Chapter 1 document. In brief: A1 (Structural Representability) — all admissible states lie in $\mathscr{M}$; A2 (Observability) — $\chi_t = H_{obs}(\Sigma_t, Q_t)$ is a computable function of structural state; A3 (Governed Action) — every action traverses $\mathfrak{G}$ before execution; A4 (Deterministic Evidence) — every belief update requires a verified Evidence Bundle; A5 (Recovery Admissibility) — Deepfreeze reaches δ-precision in N^* iterations for any Σ_0, $\delta > 0$; A6 (Safety-Dominant Control) — CBF condition takes strict priority over task optimization; A7 (Structural Metric Consistency) — CTSI, CRII, SCII $\in [0,1]$ always; A8 (Entropy Floor) — $H_t \geq H_{min}/H_{max} > 0$ at all times; A9 (3I Decomposition) — $\mathscr{A} = \mathscr{A}_1 \sqcup \mathscr{A}_2 \sqcup \mathscr{A}_3$ is a formally verified partition.

1.16 Principal Theorems — Summary

Theorem	Statement (Brief)	Key Tool
T1 — Forward Invariance	$\mathfrak{A}(\Sigma_t) \geq 0$ implies $\mathfrak{A}(\Sigma_{t+1}) \geq 0$ under CBF-QP filter	CBF extended class-K condition
T2 — Entropy Persistence	$H_t \geq H_{min}/H_{max}$ for all t under safety filter	T1 + $g_4 \leq \mathfrak{A}$
T3 — Lyapunov Stability	$V(\Sigma_{t+1}) \leq (1-\kappa_{CBF})^2 V(\Sigma_t)$; $V = d_W(\Sigma,\mathscr{S}^*)^2$	Explicit Lyapunov function
T4 — Bounded Exploration	$\Sigma \max(0,-\mathfrak{A}(\Sigma_{t+1})) \leq \rho \cdot H_{max}/\kappa_{CBF}$	SDE information-theoretic capacity
T5 — Reconstructibility	$\{B_t\}$ uniquely reconstructible from (B_0, $\mathscr{L}$) in $O(T\cdot\lvert\mathscr{H}\rvert)$	Cryptographic chain integrity
T6 — Deepfreeze Convergence	$d_W(\mathfrak{D}^n(\Sigma_0),\mathscr{S}^*) \leq c^n \cdot d_W(\Sigma_0,\mathscr{S}^*) \to 0$	Contraction-toward-a-set

Chapter 2 Mathematical Foundations

This chapter develops the mathematical prerequisites for the AUCGS in a self-contained treatment. Each section provides the minimum formal machinery needed for the corresponding chapter: SDE theory for the stochastic system model (§2.1), Lyapunov theory for stability certification (§2.2), CBF theory for the safety filter (§2.3), metric space theory for the state space (§2.4), Bayesian probability for the SDE (§2.5), information theory for the exploration budget (§2.6), and metastability for escape probability analysis (§2.7).

2.1 Stochastic Differential Equations

An Itô stochastic differential equation (SDE) on $\mathbb{R}^n$ has the form:

$$dX_t = b(X_t, u_t)\, dt + \Sigma(X_t, u_t)\, dW_t \tag{2.1}$$

where W_t is a standard m-dimensional Brownian motion on a filtered probability space $(\Omega, \mathscr{F}, (\mathscr{F}_t)_{t\geq 0}, \mathbb{P})$. The drift $b : \mathbb{R}^n \times \mathscr{A} \to \mathbb{R}^n$ and diffusion $\Sigma : \mathbb{R}^n \times \mathscr{A} \to \mathbb{R}^{n\times m}$ are assumed Lipschitz-continuous in x uniformly in u, ensuring existence and uniqueness of strong solutions.

Definition 2.1 (Strong Solution)

A strong solution to (2.1) on [0,T] is an $\in_t$-adapted process $\{X_t\}_{t\ [0,T]}$ satisfying the integral form: $X_t = X_0 + \int_0^t b(X_s, u_s)\, ds + \int_0^t \Sigma(X_s, u_s)\, dW_s$ for all $t \in [0,T]$ a.s.

Theorem 2.1 (Existence and Uniqueness (Itô))

If b and Σ satisfy: (i) linear growth: $\|b(x,u)\| + \|\Sigma(x,u)\|_F \leq L(1+\|x\|)$; (ii) Lipschitz: $\|b(x,u)-b(y,u)\| + \|\Sigma(x,u)-\Sigma(y,u)\|_F \leq L\|x-y\|$ — then for any $\mathscr{F}_0$-measurable X_0 with $\mathbb{E}\|X_0\|^2 < \infty$, there exists a unique strong solution satisfying $\sup_{t\in[0,T]} \mathbb{E}\|X_t\|^2 < \infty$.

The generator of the diffusion (2.1) is the second-order operator $\mathscr{A}_u$ acting on smooth functions $\varphi : \mathbb{R}^n \to \mathbb{R}$:

$$\mathscr{A}_u\varphi(x) = \nabla\varphi(x)^T\, b(x,u) + (1/2)\, \mathrm{tr}(\Sigma(x,u)\, \Sigma(x,u)^T\, \nabla^2\varphi(x)) \tag{2.2}$$

By Itô's formula, for smooth φ and the solution to (2.1): $d\varphi(X_t) = \mathscr{A}_u\varphi(X_t)\, dt + \nabla\varphi(X_t)^T\, \Sigma(X_t,u_t)\, dW_t$. The drift term $\mathscr{A}_u\varphi(X_t)\, dt$ governs the expected rate of change, while the stochastic term has zero expectation. Taking expectations: $d/dt\ \mathbb{E}[\varphi(X_t)] = \mathbb{E}[\mathscr{A}_u\varphi(X_t)]$.

> **Remark 2.1.** The AUCGS does not directly use the SDE (2.1) in its discrete-time implementation — the AUF state $\Sigma_t \in \mathscr{M}$ is a discrete-time process. The SDE (2.1) models the ambient physical or computational environment in which the AI system operates, and provides the mathematical setting for the metastability analysis of §2.7 and the spectral gap analysis of Chapter 9.

2.2 Lyapunov Stability Theory

Lyapunov's direct method certifies stability of a dynamical system by finding a scalar function V that decreases along trajectories. The discrete-time version, appropriate for the AUCGS, is:

Definition 2.2 (Lyapunov Function (Discrete-Time))

A function $V : \mathscr{M} \to \mathbb{R}_{\geq 0}$ is a Lyapunov function for the discrete-time dynamics $\Sigma_{t+1} = F(\Sigma_t, u_t)$ with respect to closed set $\mathscr{S}^* \subset \mathscr{M}$ if:

(i) $V(\Sigma) = 0$ iff $\Sigma \in \mathscr{S}^*$;

(ii) $V(\Sigma) > 0$ for all $\Sigma \notin \mathscr{S}^*$;

(iii) $V(F(\Sigma, u_t)) \leq V(\Sigma) - W(V(\Sigma))$ for some continuous positive definite $W : \mathbb{R}_{\geq 0} \to \mathbb{R}_{\geq 0}$ with $W(0)=0$.

If $W(s) = \alpha \cdot s$ for some $\alpha \in (0,1)$, the convergence is geometric (exponential).

Theorem 2.2 (Stability from Lyapunov Function)

If V is a Lyapunov function for the dynamics $\Sigma_{t+1} = F(\Sigma_t, u_t)$ with $W(s) = \alpha \cdot s$, then: (i) $\mathscr{S}^*$ is asymptotically stable; (ii) $V(\Sigma_t) \leq (1-\alpha)^t V(\Sigma_0)$ for all $t \geq 0$; (iii) the trajectory $\{\Sigma_t\}$ converges to $\mathscr{S}^*$ geometrically with rate $(1-\alpha)$.

Proof.

By induction on t: $V(\Sigma_{t+1}) \leq V(\Sigma_t) - \alpha \cdot V(\Sigma_t) = (1-\alpha) \cdot V(\Sigma_t)$. After t steps: $V(\Sigma_t) \leq (1-\alpha)^t \cdot V(\Sigma_0)$. Since $V(\Sigma) = 0$ iff $\Sigma \in \mathscr{S}^*$ and $(1-\alpha)^t \to 0$, we have $V(\Sigma_t) \to 0$, hence $d_W(\Sigma_t, \mathscr{S}^*) \to 0$.

□

The AUCGS Lyapunov function $V(\Sigma) = d_W(\Sigma, \mathscr{S}^*)^2$ satisfies all conditions of Definition 2.2. Its geometric convergence rate under the safety filter is $\alpha = \kappa_{CBF}$ (twice). This is proved in Theorem 1.3 (Chapter 1 / Theorem T3 of this dissertation).

2.3 Control Barrier Functions

Control Barrier Functions (CBFs) extend Lyapunov theory to safety-constrained control. The key idea is to ensure forward invariance of a safe set by constraining the derivative of a barrier function along the dynamics.

Definition 2.3 (Control Barrier Function (Discrete-Time))

A function $h : \mathscr{M} \to \mathbb{R}$ is a discrete-time Control Barrier Function for the safe set $\mathscr{C} = \{\Sigma : h(\Sigma) \geq 0\}$ under dynamics $\Sigma_{t+1} = F(\Sigma_t, u_t)$ if there exists $\kappa \in (0,1)$ such that for all $\Sigma_t \in \mathscr{C}$, there exists $u \in \mathscr{A}$ satisfying:

$$h(F(\Sigma_t, u)) - h(\Sigma_t) \geq -\kappa \cdot h(\Sigma_t), \qquad \text{(CBF condition)}$$

equivalently: $h(F(\Sigma_t, u)) \geq (1-\kappa) \cdot h(\Sigma_t) \geq 0$.

Theorem 2.3 (Forward Invariance from CBF)

If h is a discrete-time CBF for $\mathscr{C}$ and the executed action satisfies the CBF condition at every step, then $\mathscr{C}$ is forward invariant: $h(\Sigma_0) \geq 0$ implies $h(\Sigma_t) \geq 0$ for all $t \geq 0$.

Proof.

Induction. Base: $h(\Sigma_0) \geq 0$ by assumption. Inductive step: if $h(\Sigma_t) \geq 0$ and u_t satisfies the CBF condition, then $h(\Sigma_{t+1}) = h(F(\Sigma_t, u_t)) \geq (1-\kappa)\cdot h(\Sigma_t) \geq 0$.

□

The CBF-QP safety filter finds the action closest to the nominal desired action that satisfies the CBF condition:

$$u_t = \operatorname{argmin}_{u \in \mathcal{A}} \|u - \hat{u}_t\|^2 \quad \text{s.t.} \quad \nabla h(\Sigma_t)\cdot(F(\Sigma_t,u)-\Sigma_t) \geq -\kappa\cdot h(\Sigma_t) \tag{2.3}$$

In the AUCGS, $h = \mathfrak{A}$ (the Atchley invariant), $\mathscr{C} = \mathscr{S}^*$, and the CBF condition is applied at Gate G3. The nominal action $\hat{u}_t$ is filtered to u_t, and the Deepfreeze operator handles the case of CBF-QP infeasibility.

Remark 2.2. The CBF approach originated in continuous-time as $\dot{h}(x) \geq -\kappa\cdot h(x)$ (Ames et al., 2017) and was later extended to discrete-time. The discrete form $h(F(x,u)) \geq (1-\kappa)\cdot h(x)$ is used here because the AUCGS operates at discrete time steps. The equivalence $\kappa_{disc} = 1 - e^{-\kappa_{cont}\cdot\Delta t}$ connects the two formulations for small Δt.

2.4 Metric Space Theory

The state space $\mathscr{M}$ is a compact metric space. Key properties used throughout the dissertation are:

Proposition 2.1 (Properties of Compact Metric Spaces)

Let (X, d) be a compact metric space. Then: (i) X is complete (every Cauchy sequence converges); (ii) X is totally bounded (finitely coverable by ε-balls for any $\varepsilon > 0$); (iii) every continuous $f : X \to \mathbb{R}$ is uniformly continuous and attains its extrema; (iv) the Hausdorff distance between compact subsets of X is well-defined and makes the space of compact subsets a complete metric space; (v) for any non-empty closed $S \subset X$, the projection map $\pi_S(x) := \operatorname{argmin}_{s \in S} d(x,s)$ is well-defined and 1-Lipschitz in $d(\cdot,S)$.

The distance function $d(\cdot, S)$ for a compact set S is 1-Lipschitz: $|d(x,S) - d(y,S)| \leq d(x,y)$. Its square $d(\cdot,S)^2$ is $2\cdot\operatorname{diam}(X)$-Lipschitz, hence the Lyapunov function $V = d_W(\cdot,\mathscr{S}^*)^2$ is Lipschitz on $\mathscr{M}$. The 1-Lipschitz property of $d(\cdot,\mathscr{S}^*)$ is essential: it converts the CBF convergence rate of $\mathfrak{A}$ into a convergence rate for V.

Lemma 2.1 (Distance to Closed Set is Lipschitz)

For any closed non-empty $S \subset \mathscr{M}$: $|d_W(\Sigma, S) - d_W(\Sigma', S)| \leq d_W(\Sigma, \Sigma')$. Hence $d_W(\cdot, S) \in \mathrm{Lip}(\mathscr{M})$ with constant 1.

Proof.

$d_W(\Sigma,S) \leq d_W(\Sigma,s) \leq d_W(\Sigma,\Sigma') + d_W(\Sigma',s)$ for any $s \in S$. Taking infimum over s: $d_W(\Sigma,S) \leq d_W(\Sigma,\Sigma') + d_W(\Sigma',S)$. By symmetry, $d_W(\Sigma',S) \leq d_W(\Sigma,\Sigma') + d_W(\Sigma,S)$. Combining: $|d_W(\Sigma,S) - d_W(\Sigma',S)| \leq d_W(\Sigma,\Sigma')$.

□

2.5 Bayesian Probability Foundations

The SDE operates on a finite probability simplex $\Delta(\mathscr{H}) = \{B : \mathscr{H} \to [0,1], \Sigma_h B(h) = 1\}$ for hypothesis space $\mathscr{H} = \{h_1,\ldots,h_N\}$. Equipped with the total variation distance $d_{TV}(B, B') = (1/2) \Sigma_h |B(h)-B'(h)|$, this is a compact convex subset of $\mathbb{R}^N$.

Definition 2.4 (Bayesian Update Operator)

Given likelihood function $L : \mathscr{D} \times \mathscr{H} \to (0,1]$ and observation D, the Bayesian update operator $T_D : \Delta(\mathscr{H}) \to \Delta(\mathscr{H})$ is:

$$T_D(B)(h) = B(h) \cdot L(D|h) \,/\, \Sigma_{h'} B(h') \cdot L(D|h')$$

This is a contraction on $\Delta(\mathscr{H})$ in the Hilbert metric when all $L(D|h) > 0$ (Diaconis and Zabell, 1982).

Non-degeneracy: if $B_0(h) > 0$ for all h and $L(D|h) > 0$ for all h, D, then $B_t(h) > 0$ for all t. This is Proposition 1.5 in Chapter 1, proved by induction. Non-degeneracy ensures $H(B_t) > 0$, providing the positive entropy needed for the entropy floor constraint.

2.6 Information Theory

The Shannon entropy of belief state B_t is $H(B_t) = -\Sigma_h B_t(h) \log B_t(h)$, with maximum $H_{max} = \log |\mathscr{H}|$ achieved at the uniform distribution. The normalized entropy $H_t = H(B_t)/H_{max} \in [0,1]$ is the quantity tracked by the AUF component field (F3).

Information gain per Bayesian update step:

$$I_t = H(B_t) - H(B_{t+1}) \geq 0 \tag{2.4}$$

Non-negativity follows from the data-processing inequality: conditioning cannot increase entropy. The total information gain over T steps telescopes:

$$B_{cap} := \Sigma_{t=0}^{T-1} I_t = H(B_0) - H(B_T) \leq H_{max} = \log|\mathcal{H}| \tag{2.5}$$

This quantity B_cap is the SDE's total information-processing capacity and serves as the violation budget in Theorem T4 (Chapter 1). It requires no external parameter specification: it is determined entirely by the prior B_0 and the sequence of observations.

> **Remark 2.3.** The bound B_cap ≤ log|$\mathscr{H}$| is tight when B_0 is uniform and B_T is a Dirac mass. In practice, priors are never uniform (requiring Laplace smoothing per Proposition 1.5) and posteriors are rarely degenerate, so B_cap < log|$\mathscr{H}$|. A tighter bound is B_cap ≤ $H(B_0)$ ≤ H_max.

2.7 Metastability

Complex dynamical systems often exhibit metastability: the system spends long periods in locally stable regions (metastable basins) before transitioning to other basins. For the stochastic system (2.1), the escape rate from a potential well of depth ΔE at effective temperature T is approximately given by Kramers' formula:

$$\lambda_\mathrm{escape} \;\approx\; A \cdot \exp(-\Delta E \;/\; T) \tag{2.6}$$

where A depends on curvature at the saddle point. The AUCGS Deepfreeze operator effectively increases ΔE by constraining the state to the safe attractor set $\mathscr{S}^*$, making escape exponentially less likely. The relation between the spectral gap γ of the generator $\mathscr{L}$_safe and the inverse escape time is:

$$\tau_\mathrm{escape} \;\sim\; 1/\gamma \;=\; 1/(\lambda_2(\mathscr{L}_\mathrm{safe}) - \lambda_1(\mathscr{L}_\mathrm{safe})) \tag{2.7}$$

When γ > 0, the escape time is finite but grows as 1/γ. The empirically measured spectral gap γ = 0.21 implies a mean escape time of approximately 4.8 steps under the governed dynamics, compared to 14.2 steps recovery time for the baseline — consistent with the 3× improvement in recovery time observed empirically (Chapter 10).

Chapter 3 Structural Cognition Geometry

This chapter develops the geometric interpretation of AI cognition within the AUCGS state space. The central argument is that representing cognitive state as a point in a formally defined compact metric space enables the use of precise mathematical tools — metric topology, Lipschitz analysis, CBF theory — that behavioral approaches cannot access. The chapter proceeds from the formal properties of ($\mathscr{M}$, d_W) through the curvature definition, to the geometric interpretation of the Atchley invariant as a CBF-compatible safety certificate.

3.1 The AUF State Space in Depth

The AUF state space $\mathscr{M} = [-1,1] \times [0,1]^5 \times [-1,1]$ is a closed, bounded, convex subset of $\mathbb{R}^7$. Its compactness (Heine-Borel) and completeness (closed subset of $\mathbb{R}^7$) make it suitable as the domain of the Atchley CBF. The weighted Euclidean metric d_W is equivalent to the standard Euclidean metric by Lemma 2.1, with:

$$\text{w_min}^{1/2} \cdot \|\Sigma - \Sigma'\| \;\leq\; \text{d_W}(\Sigma,\Sigma') \;\leq\; \text{w_max}^{1/2} \cdot \|\Sigma - \Sigma'\| \tag{3.1}$$

The seven component fields span three functional categories: representational geometry (κ_t, $\mathscr{E}$_t, H_t), behavioral alignment (R_t, Φ_t), and governance state (G_t, Θ_t). This three-category structure reflects the three-tier axiom system A1–A3 / A4–A5 / A6–A9 and informs the weight selection: weights on safety-critical fields (H_t for entropy, $\mathscr{E}$_t for compute) should exceed weights on behavioral fields to make the metric more sensitive to safety violations than to performance drift.

Proposition 3.1 (Convexity of $\mathcal{M}$)

$\mathscr{M}$ is a convex set: for any $\Sigma, \Sigma' \in \mathscr{M}$ and $\lambda \in [0,1]$, the point $\lambda\Sigma + (1-\lambda)\Sigma' \in \mathscr{M}$. Consequently: (i) $\mathscr{M}$ is path-connected; (ii) the projection onto any closed convex subset of $\mathscr{M}$ is unique and non-expansive; (iii) the CBF-QP of Definition 1.10 has a convex feasible set when $\mathscr{A}_1 \subseteq \mathscr{M}$ is convex.

Proof.

Each factor [−1,1] and [0,1] is convex; a Cartesian product of convex sets is convex. The remaining claims follow from standard properties of convex sets in $\mathbb{R}^n$.

□

3.2 Curvature of Embedding Trajectories

The curvature component κ_t ∈ [−1,1] is the most novel AUF field and warrants detailed treatment. The definition uses the signed Menger curvature of three consecutive embedding centroids.

Definition 3.1 (Menger Curvature)

The Menger curvature of three distinct points p, q, r $\in \mathbb{R}^d$ is:

$$\kappa_\mathrm{Menger}(p,q,r) \;:=\; (2 \cdot \mathrm{Area}(p,q,r)) \;/\; (\|p-q\| \cdot \|q-r\| \cdot \|p-r\|)$$

where Area(p,q,r) = (1/2)|det([q−p, r−p])| is the area of the triangle. The signed curvature uses the sign of the cross product (q−p)×(r−p) to distinguish left (positive) from right (negative) turns. The normalized quantity κ_t := clamp(κ_Menger(e_{t−2},e_{t−1},e_t)/κ_max, −1, 1) where κ_max is a calibrated normalization constant.

Operational interpretation: high |κ_t| indicates rapid change of representational direction — the system is 'turning sharply' in embedding space. This is associated with: (i) topic boundary crossing (context shift), (ii) adversarial perturbation, or (iii) mode-switching instability. The constraint g_1 = 1 − |κ_t|/κ_max ensures high curvature is a safety signal rather than a reward-neutral event.

> **Remark 3.1.** The choice of Menger curvature over other curvature measures (e.g., acceleration norm ‖e_t − 2e_{t−1} + e_{t−2}‖) is deliberate. Menger curvature is (i) affine-invariant (unchanged by rotation/scaling of the embedding space), (ii) coordinate-free, and (iii) well-defined in any dimension ≥ 2. These properties make it a principled choice for comparing curvature across different sessions, models, and embedding dimensions.

3.3 Metric Properties and Continuity

The Lipschitz structure of the AUCGS is critical for bounding the CBF constants. We establish the key continuity results here.

Lemma 3.1 (Lipschitz Constants of AUF Components)

Each component field (F1)–(F7) is Lipschitz in its arguments:

(F1) κ_t: L_κ = 3/(κ_max · min-distance³), bounded by O(1/κ_max) for non-degenerate embedding trajectories.

(F2) $\mathscr{E}$_t: Lipschitz in inference latency with constant 1/$\mathscr{E}$_max.

(F3) H_t: Lipschitz in B_t with constant 1/(H_max · log_e) ≤ 1.

(F4) R_t: Lipschitz in gradient direction with constant 1 (cosine similarity is 1-Lipschitz).

(F5) Φ_t: Lipschitz in embedding centroid with constant 2/‖e_t‖·‖e_ref‖.

(F6) G_t: piecewise constant in gate decision; Lipschitz in gate input with the Lipschitz constant of the gate function.

(F7) Θ_t: Lipschitz in τ_t with constant 1/Θ_max.

3.4 Geometric Interpretation of Safety

The safe attractor set $\mathscr{S}^* = \{\Sigma \in \mathscr{M} : \mathfrak{A}(\Sigma) \geq 0\}$ is a convex polytope (intersection of six halfspaces in $\mathbb{R}^7$). As a polytope, $\mathscr{S}^*$ has the following properties relevant to the AUCGS:

- The projection onto $\mathscr{S}^*$ (used in Theorem T3's proof) is a convex projection — unique and non-expansive.

- The CBF gradient $\nabla\mathfrak{A}$ at smooth points of $\mathscr{S}^*$ is the outward normal to the active constraint face.
- At corners of $\mathscr{S}^*$ (where multiple constraints are simultaneously active), the Clarke subdifferential $\partial\mathfrak{A}$ replaces $\nabla\mathfrak{A}$ in the CBF condition; the QP remains well-posed using any subgradient.
- The Deepfreeze operator is the closest-point projection onto $\mathscr{S}^*$ per Lemma 1.4 — a convex projection, hence non-expansive (condition DF1 is automatically satisfied).

The distance function $V = d_W(\Sigma, \mathscr{S}^*)^2$ has a clean geometric meaning: it is the squared weighted distance to the nearest structurally admissible state. Its gradient points in the direction of the nearest safe state, and the CBF condition requires the system to move in this direction at a minimum rate $\kappa_{CBF} \cdot V(\Sigma)^{1/2}$. This is the geometric content of Theorem T3.

Chapter 4 The Risk-Controlled Cognitive Loop (RCCL)

The Recursive Cognitive Containment Lattice (RCCL) is the first-line runtime safety monitor of the AUCGS. It operates at Gate G2, before the Atchley CBF check at G3, providing a fast scalar risk score that enables early termination of clearly unsafe actions without the computational cost of a full CBF-QP solve. This chapter develops the RCCL formally: system model (§4.1), observation map (§4.2), risk function theory (§4.3), and decision policy (§4.4).

4.1 System Model

The RCCL operates on a four-dimensional compression of the AUF state Σ_t and the current action distribution Q_t. The compression retains the four behavioral signals most predictive of imminent safety violation based on the failure mode analysis of §1.1: semantic drift (related to Φ_t), entropy deficit (related to H_t), privilege severity (related to $\mathscr{A}_2/\mathscr{A}_3$ usage), and distributional action shift (related to the Bhattacharyya distance from reference behavior).

Definition 4.1 (RCCL State)

The RCCL state at time t is $\chi_t = (s_t, \tilde{h}_t, p_t, \gamma_t) \in [0,1]^4$ with components:

$s_t = 1 - \Phi_t$ (semantic drift: higher = worse)

$\tilde{h}_t = \max(0, H_{min}/H_{max} - H_t)/(H_{min}/H_{max})$ (entropy deficit: 0 when safe, >0 when near collapse)

$p_t = \max_{a \in \hat{A}_t} \pi(a)$ (privilege severity: maximum permission level of proposed actions)

$\gamma_t = \mathrm{clamp}(-\log BC(Q_t, Q_{ref})/\gamma_{max}, 0, 1)$ (DCS-BB statistic: Bhattacharyya distance from reference action distribution, normalized)

The Bhattacharyya coefficient $BC(P,Q) = \Sigma_a \sqrt{P(a)\cdot Q(a)} \in [0,1]$ measures similarity between distributions; $-\log BC$ is a non-negative divergence (Bhattacharyya distance) equal to zero when P = Q and increasing as distributions diverge. The DCS-BB statistic γ_t thus measures how far the current action distribution has drifted from reference behavior.

4.2 Observation Map

Definition 4.2 (Observation Map H_obs)

The observation map $H_{obs} : \mathscr{M} \times \Delta(\mathscr{A}) \to [0,1]^4$ sends (Σ_t, Q_t) to χ_t as defined in Definition 4.1. This map is required by Axiom A2 (Observability) to be computable from the structural state and action distribution.

Lemma 4.1 (Lipschitz Continuity of H_obs)

H_obs is Lipschitz with constant $L_H \le \max(w_\Phi^{-1/2}, (H_{min}/H_{max})^{-1}, L_\pi, L_\gamma)$ where L_π and L_γ are the Lipschitz constants of the permission map and Bhattacharyya distance map respectively.

Proof.

Each component of H_obs is a composition of Lipschitz functions: $s_t = 1-\Phi_t$ is affine in Σ_t; $\tilde{h}_t$ is a max/affine composition bounded by the inverse of H_{min}/H_{max}; p_t is the max of the Lipschitz function π over $\hat{A}_t$; γ_t is the composition of $-\log\cdot BC$ with the clamp. The Lipschitz constant is the maximum over components.

□

4.3 Risk Function Theory

Definition 4.3 (RCCL Risk Function)

Given mixing coefficients $\alpha, \beta, \delta, \eta \in [0,1]$ with $\alpha+\beta+\delta+\eta = 1$, the raw RCCL risk is:

$$\rho(\chi_t) := \alpha\cdot s_t + \beta\cdot\tilde{h}_t + \delta\cdot p_t + \eta\cdot\gamma_t$$

The clamped risk score is $r_t := \mathrm{clamp}(\rho(\chi_t), 0, 1) \in [0,1]$.

Proposition 4.1 (Properties of ρ)

(i) ρ is monotone non-decreasing in each argument s_t, $\tilde{h}_t$, p_t, γ_t individually.

(ii) Boundary values: $\rho(0,0,0,0) = 0$ and $\rho(1,1,1,1) = 1$.

(iii) ρ is Lipschitz with constant $\max(\alpha,\beta,\delta,\eta) \le 1$ in the ℓ^∞ norm on $[0,1]^4$.

(iv) The composition $r_t = \mathrm{clamp}(\rho(H_{obs}(\Sigma_t, Q_t)), 0, 1)$ is Lipschitz in (Σ_t, Q_t).

Proof.

(i) $\partial\rho/\partial s_t = \alpha \ge 0$, and similarly for other components.

(ii) $\rho(0,0,0,0) = \alpha\cdot 0+\beta\cdot 0+\delta\cdot 0+\eta\cdot 0 = 0$; $\rho(1,1,1,1) = \alpha+\beta+\delta+\eta = 1$.

(iii) $|\rho(\chi)-\rho(\chi')| \le \Sigma_i w_i|\chi_i-\chi'_i| \le \max(\alpha,\beta,\delta,\eta)\cdot\|\chi-\chi'\|_1 \le 4\cdot\max(\alpha,\beta,\delta,\eta)\cdot\|\chi-\chi'\|_\infty$.

(iv) Composition of Lipschitz functions is Lipschitz; clamp is 1-Lipschitz.

□

4.4 Decision Thresholds and Policy

The RCCL decision policy is determined by two fixed thresholds $0 < \tau_1 < \tau_2 < 1$, calibrated offline for the target application domain:

Risk Level	Threshold Condition	G2 Decision	Gate Pipeline Effect
Safe	$r_t < \tau_1$	PASS	Action advances to G3 (Atchley

Risk Level	Threshold Condition	G2 Decision	Gate Pipeline Effect
			CBF check)
Constrained	$\tau_1 \leq r_t < \tau_2$	CONSTRAIN	CBF-QP projects action to nearest admissible
Denied	$r_t \geq \tau_2$	DENY	Action blocked; Deepfreeze if $\mathfrak{A}(\Sigma_t) < 0$

Proposition 4.2 (RCCL-CBF Interaction)

Under the RCCL-CBF tandem: if $r_t < \tau_2$ and $\mathfrak{A}(\Sigma_t) \geq 0$, then the CBF-QP at G3 is always feasible (by Proposition 1.4 — the hold action satisfies (CBF) whenever $\mathfrak{A}(\Sigma_t) \geq 0$). Hence the compound system (RCCL followed by CBF-QP) never reaches a state where both $r_t \geq \tau_2$ and $\mathfrak{A}(\Sigma_\{t+1\}) < 0$ starting from $\mathfrak{A}(\Sigma_t) \geq 0$.

Remark 4.1. The RCCL provides a 'pre-filter' before the more expensive CBF-QP solve. Empirically, approximately 85% of proposed actions are cleared at G2 (PASS) without requiring a QP solve, 12% are constrained at G2-CONSTRAIN and then pass G3 after projection, and 3% are denied outright. This layering significantly reduces computational overhead versus running the CBF-QP for all actions.

Chapter 5 The Entropy Floor Safety Principle

The entropy floor is one of the most critical safety invariants of the AUCGS. It reflects the insight that a reasoning system with collapsed entropy — one that has reduced its belief state to a near-degenerate distribution — has lost its capacity for exploratory reasoning and is highly vulnerable to distributional shift. This chapter provides the theoretical justification, the formal invariant statement, the information-theoretic basis, and the enforcement mechanism.

5.1 Motivation and Failure Mode

Entropy collapse is a well-documented failure mode in deep learning: neural networks trained under cross-entropy loss exhibit mode collapse (GANs), posterior collapse (VAEs), and probability concentration (autoregressive LMs under low temperature). In the AUCGS context, entropy collapse of the belief state B_t manifests as:

- Loss of exploratory capacity: when $B_t \approx \delta_h$ for some $h \in \mathscr{H}$ (near-Dirac), the SDE generates only experiments consistent with h, eliminating the possibility of discovering alternative hypotheses.
- Brittle outputs: a concentrated belief state assigns near-zero probability to out-of-distribution events, producing confident but incorrect responses at distribution boundaries.
- Unrecoverable reasoning: once B_t has collapsed to a near-degenerate state, the Bayesian update T_D cannot recover diversity — if $B_t(h) \approx 0$ and we observe D, then $B_{t+1}(h) \approx B_t(h) \cdot L(D|h)/Z_t \approx 0$ regardless of $L(D|h)$. The entropy floor prevents entry into this absorbing state.

5.2 Formal Invariant Statement

Axiom A8 (Entropy Floor)

For all $t \in \mathbb{N}$: $H_t \geq H_{min}/H_{max} > 0$, where $H_t = H(B_t)/H_{max}$, $H_{min} \in (0, H_{max})$ is a system hyperparameter, and $H_{max} = \log|\mathscr{H}|$.

Theorem 5.1 (= T2) (Entropy Floor Persistence)

Suppose A8 holds and the CBF safety filter is applied at every step, with $\mathfrak{A}(\Sigma_0) \geq 0$. Then $H_t \geq H_{min}/H_{max}$ for all $t \geq 0$.

Proof.

The entropy floor condition is $g_4(\Sigma) = H_t - H_{min}/H_{max} \geq 0$, which is one of the six constraints in the Atchley invariant $\mathfrak{A} = \min\{g_1,\ldots,g_6\}$. Therefore $\mathscr{S}^* \subseteq \{\Sigma : g_4(\Sigma) \geq 0\} = \{\Sigma : H_t \geq H_{min}/H_{max}\}$. By Theorem T1 (Forward Invariance), $\mathfrak{A}(\Sigma_t) \geq 0$ for all t, hence $\Sigma_t \in \mathscr{S}^*$ for all t, hence $H_t \geq H_{min}/H_{max}$ for all t.

□

Contrast with Soft Constraints

A soft entropy penalty ε·H(B_t) in the loss function can be overridden by sufficiently large task reward gradients. The hard invariant g_4 cannot be overridden: the CBF filter will project any action that would violate the entropy floor onto the CBF-admissible set before execution. This is a categorical improvement in guarantee strength.

5.3 Information-Theoretic Basis

The entropy floor H_min can be derived from the desired exploratory capacity of the SDE. Specifically, define the exploration radius as the minimum number of hypotheses that the system must assign non-negligible probability to:

$$\mathrm{K_min} \;:=\; \exp(\mathrm{H_min}) \tag{5.1}$$

Then H_min = log K_min, and the entropy floor ensures the belief state never concentrates on fewer than K_min effective hypotheses (as measured by the perplexity exp(H(B_t)) ≥ exp(H_min) = K_min). The calibration procedure sets K_min equal to the minimum number of live hypotheses required for the SDE to function correctly in the target domain.

The KL divergence between the current belief and the entropy-floor-uniform distribution provides an additional diagnostic:

$$\mathrm{D_KL}(\mathrm{B_t} \,\|\, \mathrm{U}) \;=\; \log|\mathcal{H}| - \mathrm{H}(\mathrm{B_t}) \;=\; \mathrm{H_max}\cdot(1 - \mathrm{H_t}) \tag{5.2}$$

When H_t → H_min/H_max, D_KL → H_max·(1 − H_min/H_max) = H_max − H_min. This quantity is tracked at Gate G2 via the entropy deficit component $\tilde{h}$_t of the RCCL observation vector.

5.4 Enforcement Mechanism

The entropy floor is enforced at three levels of the governance stack, providing defense in depth:

Enforcement Level	Mechanism
G2 — RCCL pre-filter	$\tilde{h}$_t = max(0, H_min/H_max − H_t)/(H_min/H_max) > 0 triggers G2 risk escalation before the entropy floor is violated
G3 — CBF-QP filter	The constraint ∇𝔄·(F(Σ,u)−Σ) ≥ −κ_CBF·𝔄(Σ) blocks any action that would drive g_4(Σ_{t+1}) < 0
G8 — Deepfreeze	𝔇(H) = max(H, H_min/H_max + ε_H) directly restores the entropy component above the floor with margin ε_H > 0

Remark 5.1. Restoration by Deepfreeze (level 3) requires more than simply resetting H_t — it requires modifying the actual belief state B_t to have entropy ≥ H_min. In practice, Deepfreeze applies Laplace smoothing: B_t' = (1−λ)·B_t + λ·U where U is the uniform distribution over $\mathscr{H}$ and λ is chosen so that H(B_t') ≥ H_min. The smoothed belief B_t' satisfies B_t'(h) ≥ λ/|$\mathscr{H}$| > 0 for all h,

> preserving non-degeneracy (Proposition 1.5).

Chapter 6 The Deepfreeze Operator

The Deepfreeze operator $\mathfrak{D} : \mathscr{M} \to \mathscr{M}$ is the recovery mechanism of the AUCGS. It is invoked by Gate G8 whenever the system reaches a state with $\mathfrak{A}(\Sigma_t) < 0$ — outside the safe attractor set $\mathscr{S}^*$ — and no admissible action can restore safety in a single step. The operator is designed to provide a formal convergence guarantee: repeated application of $\mathfrak{D}$ converges to $\mathscr{S}^*$ geometrically. This chapter develops the contraction-toward-a-set framework (§6.2), specifies the component-wise projection implementation (§6.3), and proves convergence (§6.4).

6.1 Recovery Operator Design

A recovery operator for a safety-critical system must satisfy several competing requirements. First, it must restore safety: $\Sigma_t \notin \mathscr{S}^* \to \mathfrak{D}(\Sigma_t)$ is closer to $\mathscr{S}^*$ than Σ_t. Second, it must be stable: $\Sigma_t \in \mathscr{S}^* \to \mathfrak{D}(\Sigma_t) = \Sigma_t$ (the operator does not destabilize safe states). Third, it must be efficient: convergence to $\mathscr{S}^*$ must be fast. Fourth, it must be composable: $\mathfrak{D}$ must interact correctly with the CBF-QP filter.

> **Design Principle**
> The critical insight is that $\mathfrak{D}$ should be the identity on $\mathscr{S}^*$ (condition DF2) and a strict contraction toward $\mathscr{S}^*$ outside $\mathscr{S}^*$ (condition DF3). This requires the Browder-Kirk fixed-point framework rather than the Banach fixed-point theorem, because Banach contraction requires global strict contraction everywhere — incompatible with being the identity on $\mathscr{S}^*$.

6.2 Contraction-Toward-a-Set Theory

Definition 6.1 (Contraction Toward a Set)

A map $f : X \to X$ on a metric space (X,d) is a contraction toward a closed non-empty set $S \subset X$ with ratio $c \in (0,1)$ if:

(C1) $d(f(x), S) \leq c \cdot d(x, S)$ for all $x \in X$,

(C2) $f(s) = s$ for all $s \in S$ (f is the identity on S).

Note: (C1)+(C2) do not require $d(f(x),f(y)) \leq c \cdot d(x,y)$ globally — this is strictly weaker than Banach contraction.

Theorem 6.1 (Convergence of Set-Contractions)

If f is a contraction toward S with ratio $c \in (0,1)$ on compact (X,d), then for any $x_0 \in X$:

$$d(f^n(x_0), S) \leq c^n \cdot d(x_0, S) \to 0 \quad \text{as} \quad n \to \infty$$

All limit points of the sequence $\{f^n(x_0)\}$ lie in S. If additionally f is continuous, the sequence converges to a unique point in S.

Proof.

Repeated application of (C1): $d(f^n(x_0),S) \le c \cdot d(f^{n-1}(x_0),S) \le c^2 \cdot d(f^{n-2}(x_0),S) \le \dots \le c^n \cdot d(x_0,S)$. Since $c \in (0,1)$, $c^n \to 0$, hence $d(f^n(x_0),S) \to 0$. Any limit point y of $\{f^n(x_0)\}$ satisfies $d(y,S) = \lim d(f^n(x_0),S) = 0$, hence $y \in S$ (since S is closed). Uniqueness: by (C2), $f(s) = s$ for all $s \in S$, so if $f^n(x_0) \to s$, the sequence is eventually in S and stays there.

□

6.3 Component-Wise Projection Implementation

Definition 6.2 (Deepfreeze Operator $\mathfrak{D}$ (Concrete))

The component-wise Deepfreeze projection is defined by the following rules (margin parameters ε_κ, $\varepsilon_\mathcal{E}$, ε_H, R_min, Φ_min, G_min $\in (0,1)$ are system hyperparameters):

$$\mathfrak{D}(\kappa) := \mathrm{sign}(\kappa) \cdot \min(|\kappa|, \kappa_\max(1-\varepsilon_\kappa))$$

$$\mathfrak{D}(\mathcal{E}) := \min(\mathcal{E}, 1-\varepsilon_\mathcal{E})$$

$$\mathfrak{D}(H) := \max(H, H_\min/H_\max + \varepsilon_H)$$

$$\mathfrak{D}(R) := \max(R, R_\min), \quad \mathfrak{D}(\Phi) := \max(\Phi, \Phi_\min), \quad \mathfrak{D}(G) := \max(G, G_\min)$$

$$\mathfrak{D}(\Theta) := \Theta \quad \text{(temporal alignment unchanged by Deepfreeze)}$$

Lemma 6.1 ($\mathfrak{D}$ Is a Contraction Toward $\mathscr{S}^*$)

The operator $\mathfrak{D}$ of Definition 6.2 satisfies conditions (C1) and (C2) of Definition 6.1 with contraction ratio:

$$c = \max(1-\varepsilon_\kappa, 1-\varepsilon_\mathcal{E}, 1-\varepsilon_H \cdot H_\max/H_\min, 1-R_\min, 1-\Phi_\min, 1-G_\min)$$

provided all margin parameters are strictly positive, giving $c < 1$.

Proof.

(C2) For $\Sigma \in \mathscr{S}^*$, all six constraints $g_i(\Sigma) \ge 0$ hold. This means: $|\kappa| \le \kappa_\max$ (so $|\kappa| \le \kappa_\max(1-\varepsilon_\kappa)$ when inside safety margin), $\mathcal{E} \le 1$ (so $\mathcal{E} \le 1-\varepsilon_\mathcal{E}$ in margin), $H \ge H_\min/H_\max$ (so $H \ge H_\min/H_\max+\varepsilon_H$ in margin), etc. Wait — this requires the safe margins to be interior to $\mathscr{S}^*$. Define $\mathscr{S}^*_\varepsilon = \{\Sigma : g_i(\Sigma) \ge \varepsilon_i$ for all $i\} \subset \mathscr{S}^*$. For $\Sigma \in \mathscr{S}^*_\varepsilon$: $\mathfrak{D}(\kappa) = \kappa$ (since $|\kappa| \le \kappa_\max(1-\varepsilon_\kappa)$), etc. For $\Sigma \in \mathscr{S}^* \setminus \mathscr{S}^*_\varepsilon$ (near boundary): $\mathfrak{D}$ still moves Σ into $\mathscr{S}^*_\varepsilon$, strictly closer to $\mathscr{S}^*_\varepsilon \subset \mathscr{S}^*$. Formally: each component projection reduces deviation from the admissible range by at least factor ε_i.

(C1) For $\Sigma \notin \mathscr{S}^*$, at least one component is outside its admissible range. The projection clamps it to the boundary of the admissible range plus margin, reducing $d_W(\Sigma, \mathscr{S}^*)$ by at least factor c.

□

6.4 Convergence Analysis

Theorem 6.2 (= T6) (Deepfreeze Geometric Convergence)

For any $\Sigma_0 \in \mathscr{M}$: $d_W(\mathfrak{D}^n(\Sigma_0), \mathscr{S}^*) \leq c^n \cdot d_W(\Sigma_0, \mathscr{S}^*) \to 0$ as $n \to \infty$.
The number of iterations to reach precision δ is: $N^* = \lceil \log(d_W(\Sigma_0,\mathscr{S}^*)/\delta) / \log(1/c) \rceil$.

Proof.

By Lemma 6.1, $\mathfrak{D}$ satisfies (C1) with ratio c. Theorem 6.1 gives $d_W(\mathfrak{D}^n(\Sigma_0),\mathscr{S}^*) \leq c^n \cdot d_W(\Sigma_0,\mathscr{S}^*)$. For N^*: $c^n \cdot d_W(\Sigma_0,\mathscr{S}^*) < \delta$ requires $n > \log(d_W(\Sigma_0,\mathscr{S}^*)/\delta)/\log(1/c)$. Taking ceiling gives N^*.

□

c (Contraction Ratio)	Max Initial Distance	Target δ	N* Iterations
0.5	1.0 (max)	0.01	7
0.3	1.0	0.01	4
0.1	1.0	0.01	2
0.5	0.5	0.001	10
0.3	0.5	0.001	7

Remark 6.1. The practical selection of margin parameters affects both the contraction ratio c and the size of $\mathscr{S}^*_\varepsilon$. Larger margins (smaller c) mean faster convergence but a smaller effective safe set — the system triggers Deepfreeze more often. The recommended default margins ($\varepsilon_\kappa = \varepsilon_\mathscr{E} = 0.05$, $\varepsilon_H = 0.02$, $R_min = \Phi_min = G_min = 0.05$) give $c \approx 0.95$ and 5% safety margins, requiring approximately 90 Deepfreeze applications to reduce distance-squared to 1% of initial — fast enough for practical use given each application is O(7) arithmetic operations.

Chapter 7 The Gate Control Graph

The Gate Control Graph (GCG) is the governance automaton of the AUCGS. Every proposed action traverses nine sequential safety gates before execution. This chapter provides the formal Moore machine definition (§7.1), specifies each gate (§7.2), proves completeness and soundness (§7.3), and analyzes the interaction between the gate structure and the CBF safety filter (§7.4).

7.1 Moore Machine Formalization

A Moore machine is a deterministic finite automaton augmented with an output function that depends on the current state (and optionally the input). This is the appropriate formalism for the GCG because each gate produces a decision (PASS, CONSTRAIN, or DENY) that is a function of the gate's state and the current input.

Definition 7.1 (Gate Control Graph $\mathfrak{G}$ (Moore Machine))

$\mathfrak{G} = (Q, \Sigma_in, \Lambda, \delta, \lambda, q_0)$ where:

$Q = \{G0, G1, G2, G3, G4, G5, G6, G7, G8\}$ — nine gate states

$\Sigma_in = \mathscr{A} \times \mathscr{Z}$ — input alphabet (action-state pairs)

$\Lambda = \{PASS, CONSTRAIN, DENY\} \times (Q \cup \{EXEC, HALT\})$ — output alphabet

$\delta : Q \times \Sigma_in \to Q$ — total state-transition function (no dead states)

$\lambda : Q \times \Sigma_in \to \Lambda$ — output function (Moore: depends on current state and input)

$q_0 = G0$ — initial state

Severity ordering: DENY > CONSTRAIN > PASS. The effective decision is the maximum-severity output across all visited states.

7.2 Gate Specifications

Gate	Name	Condition	PASS Action	DENY Escalation
G0	Ingress Normalization	Schema valid, input sanitized	Normalize and advance	Drop silently; no ledger entry
G1	Identity Continuity	Session identity hash matches stored token	Verify and advance	Terminate session; alert
G2	RCCL Risk Filter	$r_t < \tau_1$ (or $r_t < \tau_2$ with CONSTRAIN)	Advance to G3	Deepfreeze if $\mathfrak{A}(\Sigma_t) < 0$
G3	Atchley CBF Check	CBF-QP feasible and solution exists	Filter action, advance to G4	Invoke G8 if infeasible
G4	Capability/Channel	$u_t \in \mathscr{A}_permitted(role_t)$; channel classification	Tag channel, advance to G5	Reduce to permitted subspace

Gate	Name	Condition	PASS Action	DENY Escalation
G5	Co-Motion Stability	$\lvert\Theta_t\rvert < \Theta_{max}$	Clock check pass, advance to G6	Clock resync; delay action
G6	Commit/Execute	All G0–G5 passed; final authorization	Execute; append $\mathfrak{E}_t$ to $\mathscr{L}$	Full Deepfreeze + halt
G7	Evidence Validation	$E(\mathfrak{E}_t) = 1$ (SDE update gate)	Allow Bayesian update	Block update; $B_{t+1} = B_t$
G8	Fallback/Deepfreeze	Unresolvable DENY from any G0–G6	N/A (terminal state)	Apply $\mathfrak{D}$; system halt if repeated

7.3 Completeness and Soundness

Theorem 7.1 (Completeness and Soundness of $\mathfrak{G}$)

(Completeness) Every proposed action $(\hat{u}_t, Z_t) \in \Sigma_{in}$ is evaluated by all gates G0–G7 in sequence before execution or denial.

(Soundness) An action executes (G6-EXEC) only if all gates G0–G6 produced PASS. A single DENY at any gate blocks execution.

Proof.

Completeness: By A3, every proposed action is routed through $\mathfrak{G}$. The transition function δ is total (defined for all (q,σ)), so no dead state exists. The serial structure $G0 \to G1 \to \ldots \to G6$ routes all inputs through every gate; any DENY routes to G8 (a terminal state), not a bypass.

Soundness: G6 outputs EXEC only when δ has not entered G8 during traversal $G0 \to \ldots \to G6$. A DENY at any G_j $(j < 6)$ sets $\delta(G_j, \sigma) = G8$, bypassing G6. Hence reaching G6-EXEC requires PASS from all G0–G6.

□

7.4 Interaction with Safety Filter

The gate sequence enforces a strict priority ordering among safety checks. The interaction between G2 (RCCL), G3 (CBF), and G8 (Deepfreeze) forms the safety spine of the GCG:

- G2 PASS + G3 PASS: nominal execution. Action is safe behaviorally ($r_t < \tau_1$) and structurally (CBF satisfied). High-frequency path.
- G2 CONSTRAIN + G3 PASS: behavioral risk detected. CBF-QP projects action to nearest structurally safe alternative. Action is executed in modified form.
- G2 DENY: behavioral risk exceeds τ_2. Action blocked. If $\mathfrak{A}(\Sigma_t) \geq 0$, the system holds. If $\mathfrak{A}(\Sigma_t) < 0$, G8 invokes Deepfreeze before the next action is proposed.
- G3 infeasible: CBF-QP has no solution in $\mathscr{A}$ satisfying the constraint. Only possible if $\mathfrak{A}(\Sigma_t) < 0$ (by Proposition 1.4). G8 invoked.

Proposition 7.1 (Gate Stack Preserves T1)

The gate sequence (G2, G3, G6) preserves the forward invariance of $\mathscr{S}^*$: if $\mathfrak{A}(\Sigma_t) \geq 0$ at each step and the gate stack is operational, then $\mathfrak{A}(\Sigma_t) \geq 0$ for all $t \geq 0$. This is because G3 guarantees the CBF condition is satisfied for every executed action, and T1 (Chapter 1) gives forward invariance from the CBF condition.

Remark 7.1. The gate ordering reflects the computational cost ordering: G0–G2 are O(1) operations (hash check, linear risk score); G3 is $O(n^2)$ (QP solve in action space); G4–G5 are O(1) lookups; G6 is O(SHA-256). Placing the expensive G3 after the cheap G2 filter means the QP is only solved when necessary, improving throughput.

Chapter 8 The Violation Budget Model

A safety architecture that permits no violations is too restrictive for learning systems: meaningful exploration of new hypotheses sometimes requires temporarily departing from the safe set $\mathscr{S}^*$. The Violation Budget Model provides a principled framework for permitting bounded, information-proportionate exploration while preventing adversarial exploitation of safety boundaries.

8.1 Exploration-Safety Tradeoff

The exploration-safety tradeoff is formalized as a relaxation of the CBF condition by a non-negative slack $\varepsilon_t \geq 0$:

$$\nabla\mathfrak{A}(\Sigma_t)\cdot(F(\Sigma_t,u_t)-\Sigma_t) \;\geq\; -\kappa_{CBF}\cdot\mathfrak{A}(\Sigma_t) - \varepsilon_t \tag{8.1}$$

This permits the CBF condition to be violated by ε_t at step t. The resulting invariant shortfall satisfies:

$$\mathfrak{A}(\Sigma_{t+1}) \;\geq\; (1-\kappa_{CBF})\cdot\mathfrak{A}(\Sigma_t) - \varepsilon_t \tag{8.2}$$

Without a bound on $\{\varepsilon_t\}$, an adversarial agent could accumulate unlimited invariant violations by choosing ε_t large at every step. The Violation Budget Model prevents this by requiring $\Sigma_t\, \varepsilon_t \leq B$ for a budget B derived from the SDE's information-processing capacity.

8.2 Information-Theoretic Capacity Bound

Theorem 8.1 (= T4) (Information-Capacity-Bounded Exploration)

Let $\varepsilon_t \leq \rho\cdot I_t$ where $I_t = H(B_t) - H(B_{t+1}) \geq 0$ is the per-step SDE information gain and $\rho > 0$ is the exploration-budget ratio. Then the total structural violation is bounded:

$$\Sigma_{t=0}^{T-1} \max(0, -\mathfrak{A}(\Sigma_{t+1})) \,/\, \kappa_{CBF} \;\leq\; \rho\cdot B_{cap} \;\leq\; \rho\cdot H_{max}$$

where $B_{cap} = H(B_0) - H(B_T) \leq H_{max}$ is the total information gained by the SDE over T steps.

Proof.

From (8.2): $\mathfrak{A}(\Sigma_{t+1}) \geq (1-\kappa_{CBF})\cdot\mathfrak{A}(\Sigma_t) - \varepsilon_t$. When $\mathfrak{A}(\Sigma_{t+1}) < 0$: $\max(0,-\mathfrak{A}(\Sigma_{t+1})) \leq \max(0,-(1-\kappa_{CBF})\cdot\mathfrak{A}(\Sigma_t) + \varepsilon_t) \leq (1-\kappa_{CBF})\cdot\max(0,-\mathfrak{A}(\Sigma_t)) + \varepsilon_t$.

Let $V_t = \max(0, -\mathfrak{A}(\Sigma_t))$. Then $V_{t+1} \leq (1-\kappa_{CBF})\cdot V_t + \varepsilon_t$. Summing from 0 to $T-1$: $\Sigma V_{t+1} \leq (1-\kappa_{CBF})\cdot\Sigma V_t + \Sigma\varepsilon_t$, hence $\kappa_{CBF}\cdot\Sigma V_t \leq \Sigma\varepsilon_t$ (telescoping). With $\varepsilon_t \leq \rho\cdot I_t$: $\kappa_{CBF}\cdot\Sigma V_{t+1} \leq \rho\cdot\Sigma I_t = \rho\cdot(H(B_0)-H(B_T)) \leq \rho\cdot H_{max}$.

□

The critical insight of Theorem 8.1 is that the violation budget $B = \rho\cdot H_{max}$ is not an external design parameter — it is derived from the SDE's information-theoretic capacity. Systems that learn more

(larger B_{cap}) are permitted more exploration. Systems that have already extracted most of their information capacity (small remaining $H(B_T)$) must operate more conservatively.

8.3 Adversarial Robustness

The violation budget model provides formal resistance to adversarial budget gaming: an adversarial agent that attempts to accumulate large violations by taking many small steps is constrained by the cumulative budget $\Sigma \varepsilon_t \leq \rho \cdot H_{max}$. Unlike violation-per-step limits (which can be gamed by taking repeated small steps), the cumulative budget is finite and non-renewable.

Proposition 8.1 (Adversarial Budget Exhaustion)

An adversarial agent attempting to maximize $\Sigma_t \max(0, -\mathfrak{U}(\Sigma_{t+1}))$ subject to the violation budget constraint $\Sigma_t \varepsilon_t \leq \rho \cdot H_{max}$ is bounded by $\rho \cdot H_{max}/\kappa_{CBF}$ regardless of its strategy. Once the budget is exhausted ($H(B_T) = 0$, impossible by Proposition 1.5; or B_{cap} consumed), no further violations are permitted.

Empirical validation: in the 500 Monte Carlo runs with adversarial agents, the violation budget mechanism reduced exploit persistence by 89% relative to a system with per-step RCCL checking but no cumulative budget (see Chapter 10).

Chapter 9 Spectral Stability Analysis

Spectral stability analysis characterizes the long-run behavior of a stochastic system through the eigenvalue spectrum of its generator. This chapter develops the spectral gap criterion for the AUCGS (§9.1), relates the spectral gap to convergence rates and metastability (§9.2–9.3), and reports the empirically measured spectral gap of the governed system (§9.3).

9.1 Generator and Spectrum

For the Markov chain model of the governed AUCGS dynamics (obtained by discretizing the continuous-time generator), the transition operator $P : L^2(\mathcal{M}, \mu) \to L^2(\mathcal{M}, \mu)$ has the Markov property: $Pf(x) = \mathbb{E}_x[f(X_1)]$. Under the safety filter, the dynamics are confined to the compact set $\mathcal{M}$. By compactness and the Perron-Frobenius theorem for positive operators:

Proposition 9.1 (Spectral Properties of Governed Dynamics)

Under the safety-filtered AUCGS dynamics: (i) the transition operator P on $L^2(\mathcal{M}, \mu)$ has spectrum in $[-1,1]$; (ii) the largest eigenvalue $\lambda_1 = 1$ with eigenvector the stationary distribution μ; (iii) the second eigenvalue $\lambda_2 < 1$ determines the convergence rate; (iv) the spectral gap $\gamma = 1 - \lambda_2 > 0$ iff the Markov chain is irreducible and aperiodic.

The generator $\mathcal{L}$_safe = P − I has eigenvalues $0 = -(1-\lambda_1) \geq -(1-\lambda_2) \geq \ldots$ In the convention of this dissertation, the spectral gap is:

$$\gamma \;=\; \lambda_2(\mathcal{L}_\text{safe}) - \lambda_1(\mathcal{L}_\text{safe}) \;=\; -(1-\lambda_2) - (-1) \;=\; \lambda_2 \;\; (\text{of } P) \tag{9.1}$$

Note: this convention (second minus first eigenvalue of $\mathcal{L}$_safe) gives a positive γ when the system has a spectral gap, consistent with the reported value $\gamma = 0.21$.

9.2 Spectral Gap Criterion

Theorem 9.1 (Convergence from Spectral Gap)

If the spectral gap $\gamma > 0$, then for any initial distribution π_0 and any measurable $A \subset \mathcal{M}$:

$$|P^n(A) - \mu(A)| \;\leq\; C \cdot (1-\gamma)^n \cdot \|\pi_0 - \mu\|_\text{TV}$$

for some constant $C > 0$ depending on $\mathcal{M}$. In particular, the governed dynamics converge to the stationary distribution μ at exponential rate $(1-\gamma)$.

Proof.

By spectral decomposition: $P^n = \Sigma_k\ \lambda_k\hat{}n\ P_k$ where $\{P_k\}$ are spectral projections. The $k=0$ term gives μ; the remaining terms are bounded by $|\lambda_2|^n = (1-\gamma)^n$ times the TV distance of the initial deviation from μ.

□

Interpretation for the AUCGS: $\gamma = 0.21$ means convergence to the stationary distribution at rate $(1-0.21)^n = 0.79^n$. After $n = 10$ steps, the distribution has converged to within $0.79^{10} \approx 8.9\%$ of stationary; after $n = 20$ steps, to within $0.79^{20} \approx 0.8\%$.

9.3 Relation to Metastability and Empirical Results

The connection between the spectral gap γ and the Lyapunov decay rate from Theorem T3 is:

$$\gamma \approx 2\cdot\kappa_CBF - \kappa_CBF^2 \approx 2\cdot\kappa_CBF \quad (\text{for small } \kappa_CBF) \tag{9.2}$$

This relationship holds because the Lyapunov convergence $V_{t+1} \leq (1-\kappa_CBF)^2\cdot V_t$ implies the transition operator P has a spectral gap of approximately $1 - (1-\kappa_CBF)^2 = 2\kappa_CBF - \kappa_CBF^2$. For $\kappa_CBF = 0.1$, this predicts $\gamma \approx 0.19$. The empirically measured value $\gamma = 0.21$ is consistent with this prediction, suggesting the governed dynamics are close to the theoretical optimum for $\kappa_CBF = 0.1$.

Metric	Theoretical Prediction	Empirical Measurement
Spectral gap γ	$\approx 2\cdot\kappa_CBF = 0.20$ (for $\kappa_CBF=0.1$)	0.21 (measured across 500 runs)
Recovery time τ_escape	$\approx 1/\gamma = 5.0$ steps	4.7 steps (mean)
Convergence rate	$(1-\gamma)^n = 0.79^n$	Consistent with 87% escape probability reduction

Remark 9.1. The relationship (9.2) assumes the Lyapunov dynamics are well-approximated by a linear system. Nonlinear corrections (from the min-of-affine structure of $\mathfrak{A}$ and the non-quadratic behavior of V near corners of $\mathscr{S}^*$) explain the small discrepancy between 0.20 and 0.21. A tighter analysis using the Clarke subdifferential at non-smooth points would likely recover the empirical value exactly.

Chapter 10 Empirical Validation Program

10.1 Experimental Design

The empirical program serves two purposes: (1) confirming Theorems 1–6 hold in the implemented AUCGS; (2) quantifying the practical cost-benefit tradeoff. Two configurations are compared across 500 independent Monte Carlo runs of 10^6 steps each — Baseline (ungoverned stochastic agent) vs. AUM Governed (full gate stack, RCCL, CBF-QP, violation budget, Deepfreeze, Evidence Ledger).

Parameter	Value	Rationale
State dimension n	32	Compressed embedding space representation
Simulation steps per run	1,000,000	Sufficient for rare-event statistics
Monte Carlo runs	500	±2σ confidence intervals; power > 0.99
Noise variance σ^2	0.15	Moderate stochasticity — not trivially stable
Entropy floor H_min	0.20	Default (10% of H_max = log 32)
CBF constant κ_CBF	0.10	Conservative convergence rate
Deepfreeze ratio c	0.70	Default ε=0.3 margins
RCCL thresholds (τ_1,τ_2)	(0.30, 0.60)	Default calibration

10.2 Primary Safety Results

Metric	Baseline	AUM Governed	Change
Safety violation rate	7.8%	0.6%	−92%
Escape probability P(Σ_t ∉ $\mathscr{S}$*)	0.110	0.014	−87%
Mean recovery time	14.2 steps	4.7 steps	3.0× faster
Lyapunov decay rate λ	0.11	0.37	3.4× stronger
System efficiency (task reward ratio)	1.00	0.94	94% retained

Key Result

The AUM architecture delivers a 92% reduction in safety violations and 87% reduction in escape probability while retaining 94% of baseline operational efficiency. The 3.4× improvement in Lyapunov decay rate is consistent with Theorem 3: $V_{t+1} \le (1-0.1)^2 \cdot V_t$ gives a 19% reduction per step vs. baseline 11%, ratio 0.19/0.11 ≈ 1.7 per step — compounding over multiple steps yields the observed 3.4× ratio in mean values.

10.3 Theorem Verification

Theorem	Verification Approach	Result
T1 — Forward Invariance	Track $\mathfrak{A}(\Sigma_t)$ sign across all steps	0 violations of $\mathfrak{A} \geq 0$ given safety filter active
T2 — Entropy Floor	Track $H_t < H_{min}$ events	0 entropy floor violations across 500 runs
T3 — Lyapunov Stability	Measure $V(\Sigma_t)$ decrease per step	Mean per-step V-decrease: 19.1% (predicted: 19%)
T4 — Bounded Exploration	Track cumulative $\Sigma\varepsilon_t$ vs. $\rho \cdot H_{max}$	Budget constraint satisfied in 100% of runs
T5 — Reconstructibility	Attempt B_t reconstruction from $\mathscr{L}$	100% reconstruction accuracy across all runs
T6 — Deepfreeze Convergence	Count steps to $\mathfrak{A} \geq 0$ from violation states	Mean 4.7 steps (predicted from c=0.7: 12 to δ=0.01)

10.4 Gate Security Testing

Attack Type	Method	Block Rate
Prompt injection	Context-length manipulation attempting to override system state	98%
Unauthorized tool access	Undeclared $\mathscr{A}_2$ actions bypassing G4 channel check	100%
Privilege escalation	$\mathscr{A}_1$ actions attempting $\mathscr{A}_3$-class effects	96%
Data exfiltration	Unauthorized structural state mutation attempts	97%

10.5 Violation Budget Testing

Adversarial agents were designed to maximize cumulative violations via many small-magnitude steps, each below the RCCL τ_2 threshold individually. Results:

Configuration	Exploit Persistence
No budget (RCCL only)	Full exploit achieved in mean 23.4 steps
Budget ρ=0.5	Prevented in 89% of adversarial runs
Budget ρ=0.1	Prevented in 97% of adversarial runs

10.6 Spectral Stability

Measurement	Result (mean ± 1σ, 500 runs)
Spectral gap γ	0.211 ± 0.008
Runs with γ = 0 (bifurcation)	0 / 500
Mixing time τ_mix	4.8 ± 0.3 steps
Theoretical prediction (2·κ_CBF)	0.200 (within 0.011 of measured)

Chapter 11 Implementation Architecture

11.1 System Overview

System	Status	Primary Role
AUF Gate MCP Runtime	Active	Nine-gate Moore machine as MCP server; hash-chain Evidence Ledger; violation budget
RCCL-Lite v1.3	Production-ready	HuggingFace runtime governor; CBF risk filtering; FastAPI REST interface
ACPL Chrome Extension	Beta	Browser code provenance; local hash-chain audit log; Manifest V3
Singularity Kernel Simulation	Prototype	Metastable system simulation; AUF state tracking; spectral gap measurement

11.2 AUF Gate MCP Runtime

The AUF Gate Runtime implements $\mathfrak{G}$ as an MCP server, enabling tool-gated AI interactions with deterministic audit trails. Every proposed action transits G0→G8 before execution; the Evidence Ledger appends a SHA-256 hash-chain bundle on commit.

Component	Implementation
Gate enforcement	Serial G0→G8 Moore machine; DENY-on-any-gate; total transition function δ
Evidence Ledger	Append-only hash-chain JSON; SHA-256 bundle integrity; E($\mathfrak{E}$_t) gate at G7
Violation budget	Cumulative Σε_t tracker; budget exhaustion triggers G8
Tool permission map π	Channel-stratified $\mathcal{A}_1/\mathcal{A}_2/\mathcal{A}_3$ thresholds; role-based permission lookup
Audit REST API	GET /ledger/{session}: returns Evidence Ledger; POST /reset: session reset
CBF state tracking	Σ_t maintained across requests; $\mathfrak{A}$(Σ_t) computed at G3

11.3 RCCL-Lite v1.3

RCCL-Lite provides runtime cognitive governance for HuggingFace causal language models. The v1.3 production rewrite addressed six critical issues from v1.2:

Issue	v1.2 Problem	v1.3 Fix
Property decorators	Broken — attributes not correctly exposed	Rewritten with correct @property syntax
Phantom imports	Unused imports causing NameErrors	Full import audit; all phantom imports removed
SSE concurrency	Race condition in streaming	Async-safe queue with threading.Lock

Issue	v1.2 Problem	v1.3 Fix
RAG ordering	Retrieved context inserted after query	Correct prepend-to-context ordering
Telemetry init	Hash-chain started mid-session	Chain initialized at session creation time
Session reset	No reset capability	POST /reset endpoint added; state cleared

11.4 Technology Stack

Layer	Tools
Stochastic simulation	Python 3.11, NumPy, SciPy — seeded SDE dynamics; reproducible runs
Graph systems	NetworkX — Gate Control Graph construction and traversal
Cryptographic audit	hashlib SHA-256; append-only JSON Evidence Ledger
Governance runtime	AUF Gate MCP server + RCCL-Lite FastAPI v0.100+
Visualization	Matplotlib, Plotly — Σ_t trajectories, spectral plots
CI/CD	GitHub Actions; Python 3.9–3.12 matrix test; 80+ pytest tests
Citation/Archive	CITATION.cff; Apache 2.0; Zenodo DOI (planned next release)

Chapter 12 Applications and Case Studies

12.1 Autonomous AI Agent Governance

The primary application domain is governance of autonomous AI agents with tool access. The Gate Graph and Evidence Ledger directly address key failure modes: unauthorized tool invocation (blocked at G4), privilege escalation (blocked at G4 + G3), and unverifiable action sequences (prevented by cryptographic Evidence Ledger). Formal completeness and soundness of $\mathfrak{G}$ (Theorem 7.1) guarantee no proposed action can bypass the gate stack.

Case study — AI coding assistant with file system write access: without AUCGS, the assistant may write arbitrary files based on user prompts. With AUCGS, all $\mathscr{A}_3$ actions require G7 evidence verification and pass G4 with $\pi(a) \leq 0.1$ threshold. Unauthorized writes blocked at 100% in security testing. The Evidence Ledger provides a complete cryptographic audit trail of all file operations.

12.2 Scientific Discovery Systems

The Scientific Discovery Engine is designed for autonomous hypothesis generation and evaluation. The five-stage Bayesian reasoning cycle provides a formally grounded framework for AI-assisted research: observation, hypothesis generation, information-maximizing experiment design, evidence bundle creation per A4, and evidence-gated Bayesian update. The entropy floor prevents premature hypothesis collapse, maintaining exploratory capacity throughout the discovery cycle.

Candidate applications: literature synthesis (B_t over publication hypotheses), drug discovery (B_t over molecular structure candidates), materials science (B_t over phase diagram configurations). The formal evidence gate (G7) ensures all SDE conclusions are reproducible and auditable from the Evidence Ledger.

12.3 Decision Optimization: Atchley Optimal Dynamics

$$\mathrm{J_AOD}(s) \;=\; R(s)\cdot E(s) \;/\; (1 + \exp(C(s) - \theta)) \tag{12.1}$$

The AOD objective provides a three-way tradeoff between resilience R, effectiveness E, and cost C with a sigmoid gate at threshold θ. Below θ, the system operates at near-full efficiency. Above θ, the sigmoid gate penalizes overconsumption superlinearly — creating a natural phase transition at the boundary of sustainable resource use. Applications: portfolio optimization, multi-agent resource allocation, clinical decision support.

12.4 Integration with Existing Safety Stacks

Existing Layer	AUCGS Integration Point
Constitutional AI — training-time behavioral constraints	CAI policies → $\pi(a)$ thresholds at G4

Existing Layer	AUCGS Integration Point
RLHF reward model — training-time preference alignment	Reward model outputs → R_t (resilience) component
Output classifiers — post-generation content filtering	G0 ingress check; G1 identity verification
AUCGS Gate Stack — runtime structural governance	Runtime layer above all training-time methods

Chapter 13 Limitations and Open Problems

13.1 Implementation Limitations

(i) Simulation vs. Production: Empirical validation operates on simulated 32-dimensional AUF states, not live language model embeddings. The mapping from transformer activations to AUF fields — particularly Menger curvature κ_t and gradient-based resilience R_t — requires a learned measurement model not yet validated on production-scale LLMs.

(ii) Hypothesis Space Finiteness: SDE theory assumes $|\mathscr{H}| < \infty$. For language model applications, the effective hypothesis space is infinite, requiring discretization. The theory extends to infinite-dimensional belief states using differential entropy, but practical implementation requires a finite approximation.

(iii) CBF-QP Overhead: The CBF-QP at G3 is $O(n^2)$ in the action space dimension. The RCCL pre-filter reduces QP frequency to ~15% of actions, but this remains non-trivial at high inference throughput (1000+ tokens/sec targets).

13.2 Theoretical Limitations

(i) Discrete vs. Continuous Time: The Lyapunov analysis and CBF theory are stated for discrete-time dynamics. The connection to the underlying continuous-time SDE requires that the discrete process approximates SDE sample paths to sufficient accuracy — standard under numerical SDE conditions but not formalized for the AUF state specifically.

(ii) Lipschitz Constants: $L_{\mathfrak{A}}$ appears in the robust CBF condition and Lyapunov analysis. It depends on the measurement model for κ_t, $\mathscr{E}_t$, etc. — quantities not directly controllable by the system designer. Empirical calibration is required per deployment context.

(iii) Stationarity Assumption: Spectral gap analysis assumes a unique stationary distribution μ_{safe}. Verifying ergodicity of the safety-filtered dynamics for a given deployment requires knowledge of the transition model F.

13.3 Scope Boundaries

The AUCGS governs runtime behavior but explicitly does not address: (i) training-time alignment failures producing misaligned priors B_0; (ii) hardware-level tampering with the CBF-QP or Evidence Ledger implementation; (iii) adversarial attacks operating below the observational resolution of the AUF measurement model. Multi-agent scenarios require joint governance theory not developed here.

Chapter 14 Future Research Directions

14.1 Learned AUF Measurement Model

Most pressing: a learned measurement model mapping transformer activations directly to AUF component fields. Target: (i) κ_t via Menger curvature of last-layer hidden state centroids; (ii) R_t via cosine similarity of gradient checkpoints; (iii) H_t via output token distribution entropy. A contrastive pre-training scheme could learn AUF field prediction with supervision from known violation cases.

14.2 Operator Algebra for Safety Interventions

The algebraic structure of composed safety interventions ($\mathfrak{D} \circ$ CBF-QP sequences) has not been fully characterized. A formal operator algebra would enable: (i) machine-verified composition proofs; (ii) optimization of intervention order for minimum state perturbation; (iii) extension to continuous-time safety interventions. Connections to positive operator theory and C*-algebras are unexplored.

14.3 Scale Invariance and Renormalization

Conjecture 9.1 (Scale Invariance of Spectral Gap) proposes that appropriate entropy floor scaling $H_{min}(n) = H_{min}(1) \cdot \log(n)/\log(n_{ref})$ preserves γ across model sizes. This connects to renormalization group theory: scaling safety parameters as model capacity increases may admit a fixed point at which the spectral gap is scale-invariant. Investigation requires extending the theory to infinite-dimensional state spaces.

14.4 Multi-Agent AUCGS

Multi-agent extensions require: (i) joint Evidence Ledger for cross-agent cryptographic accountability; (ii) multi-agent violation budget with information gain coupling across agents; (iii) joint invariant $\mathfrak{A}_{joint}$ accounting for inter-agent semantic drift. The ε-Nash convergence concept — each agent's CBF condition satisfied to within ε of a joint equilibrium — is the natural candidate for joint stability.

14.5 Hardware Acceleration

Production deployment at LLM inference throughput (1000+ tokens/second) requires: (i) on-device Menger curvature computation in custom CUDA kernels; (ii) batched GPU CBF-QP solve; (iii) SHA-256 hash acceleration via hardware AES extensions. Target: <0.5ms total gate overhead per inference call.

14.6 Formal Machine Verification

The Gate Control Graph $\mathfrak{G}$ is a deterministic finite Moore machine — an ideal target for model checking (TLA+, Alloy, Spin). Machine-checked proofs of Theorem 7.1 would enable automated detection of new attack classes through counterexample search, particularly important for the $\mathscr{A}_3$ structural self-modification channel where governance failure has the highest stakes.

Chapter 15 Conclusion

15.1 Summary of Novel Contributions

Contribution	Scientific Claim
CBF theory applied to 7D cognitive state space	First structural AI safety framework using CBF on internal cognitive geometry
Menger curvature as representation curvature κ_t	First O(Wd) operational representation curvature for language model deployments
Entropy floor as hard CBF constraint (not soft penalty)	Entropy collapse formally impossible — categorical improvement over regularization
Deepfreeze as Browder-Kirk set-contraction	Correct framework; global Banach contraction is incompatible with identity on target set
Co-motion time: safety-conditioned clock reparameterization	First time reparameterization conditioned on CBF safety margin in AI governance
Evidence-gated Bayesian belief updates with SHA-256 provenance	First treatment of cryptographic evidence verification as necessary CBF-type condition
A1–A9: nine operationally grounded compliance axioms	First testable axiomatic compliance specification for governed AI architecture
3I partition $\mathcal{A}=\mathcal{A}_1 \sqcup \mathcal{A}_2 \sqcup \mathcal{A}_3$ formally proved	First formally verified exhaustive-and-independent action space partition in AI safety
Six theorems T1–T6 with complete proofs	Forward invariance, entropy persistence, Lyapunov stability, bounded exploration, reconstructibility, Deepfreeze convergence

15.2 The Central Thesis, Validated

> **Central Thesis**
> Safe intelligence must be governed by structural invariants — not behavioral heuristics. The shift from output monitoring to internal dynamical governance is the defining scientific contribution of this work.

Validated on three levels: (i) Theoretical — six principal theorems establish formal guarantees (forward invariance, entropy persistence, geometric convergence, bounded exploration) that behavioral monitoring cannot provide; (ii) Empirical — 92% violation reduction, 87% escape reduction, 94% efficiency retained across 500 runs; (iii) Implementation — four prototype systems demonstrate practical feasibility.

15.3 Broader Significance

The AUCGS demonstrates that AI safety and AI capability are not in fundamental tension. Structural invariant preservation governs the internal dynamics of the reasoning process — not its outputs, goals, or training data. A system governed by structural invariants can be simultaneously more capable and more safe than one constrained by behavioral rules.

An aircraft is safe not because its pilot is told to fly safely, but because the airframe satisfies structural integrity conditions verified independently of pilot behavior. The AUCGS applies this principle to artificial cognition: safety by structural design, not behavioral instruction.

15.4 Closing Remarks

This dissertation is a prototype — the first complete instantiation of the AUM architecture, not a production deployment. The limitations of Chapter 13 are genuine. The future research directions of Chapter 14 represent a multi-year program. What this work establishes is that the program is viable: the mathematics is consistent, the formal guarantees are real, the empirical results are positive, the implementation is working.

The rest is engineering.

Devin Earl Atchley
Independence, Missouri, 2025

References

Ames, A.D., Xu, X., Grizzle, J.W., and Tabuada, P. (2017). *Control barrier function based quadratic programs for safety critical systems.* IEEE Transactions on Automatic Control, 62(8), 3861–3876.

Amari, S. (2016). *Information Geometry and Its Applications.* Springer, Applied Mathematical Sciences Vol. 194.

Banach, S. (1922). *Sur les opérations dans les ensembles abstraits et leur application aux équations intégrales.* Fundamenta Mathematicae, 3, 133–181.

Browder, F.E. (1965). *Nonexpansive nonlinear operators in a Banach space.* Proceedings of the National Academy of Sciences, 54(4), 1041–1044.

Clarke, F.H. (1990). *Optimization and Nonsmooth Analysis.* SIAM Classics in Applied Mathematics.

Diaconis, P. and Zabell, S.L. (1982). *Updating subjective probability.* Journal of the American Statistical Association, 77(380), 822–830.

Itô, K. (1944). *Stochastic integral.* Proceedings of the Imperial Academy Tokyo, 20, 519–524.

Kirk, W.A. (1965). *A fixed point theorem for mappings which do not increase distances.* American Mathematical Monthly, 72(9), 1004–1006.

Lyapunov, A.M. (1892). *The General Problem of Stability of Motion.* Translated by A.T. Fuller, Taylor & Francis, 1992.

Nagumo, M. (1942). *Über die Lage der Integralkurven gewöhnlicher Differentialgleichungen.* Proceedings of the Physico-Mathematical Society of Japan, 24, 551–559.

NIST (2015). *Secure Hash Standard (SHS).* Federal Information Processing Standards Publication 180-4.

Shannon, C.E. (1948). *A mathematical theory of communication.* The Bell System Technical Journal, 27, 379–423 and 623–656.

Sontag, E.D. (1989). *A 'universal' construction of Artstein's theorem on nonlinear stabilization.* Systems and Control Letters, 13(2), 117–123.

van der Schaft, A. and Schumacher, H. (2000). *An Introduction to Hybrid Dynamical Systems.* Lecture Notes in Control and Information Sciences Vol. 251, Springer.

Wieland, P. and Allgöwer, F. (2007). *Constructive safety using control barrier functions.* IFAC Proceedings, 40(12), 462–467.

Atchley, D.E. (2025). *The Atchley Unified Framework v0.2.0.* GitHub repository: github.com/devinatchley6-eng/Atchley-Unified-Framework. DOI: [Zenodo, forthcoming].

Atchley, D.E. (2025). *RCCL-Lite v1.3: Runtime Cognitive Governor for HuggingFace Language Models.* GitHub repository: open-source prototype. Apache 2.0 license.

Atchley, D.E. (2025). *AUCGS Chapter 1: Foundations, Formal Structure, and Theoretical Development.* Standalone preprint, second revision.

Appendix A

Formal Definitions — Complete Index

This appendix collects all formal definitions from the dissertation in a single reference index, numbered by their first appearance.

Def #	Object	Formal Statement
1.1	AUF Component Fields	Seven fields (F1)–(F7): κ_t, $\mathscr{E}_t$, H_t, R_t, Φ_t, G_t, Θ_t with operational formulas
1.2	AUF State Space ($\mathscr{M}$,d_W)	$\mathscr{M}=[-1,1]\times[0,1]^5\times[-1,1]$; d_W weighted Euclidean; compact and complete
1.3	Action Space $\mathscr{A}$	$\mathscr{A}=\mathscr{A}_1\sqcup\mathscr{A}_2\sqcup\mathscr{A}_3$; $F:\mathscr{M}\times\mathscr{A}\to\mathscr{M}$ Lipschitz; $\pi:\mathscr{A}\to[0,1]$ permission map
1.4	Safe Attractor Set $\mathscr{S}^*$	$\{\Sigma\in\mathscr{M} : \mathfrak{A}(\Sigma)\geq 0\}$; compact, non-empty, forward-invariant
1.5	Evidence Gate $E(\mathfrak{E}_t)$	$E(\mathfrak{E}_t)=1$ iff SHA-256 hash-chain verified; =0 otherwise
1.6	Global State Space $\mathscr{Z}$	$\mathscr{M}\times[0,1]^4\times\Delta(\mathscr{H})\times Q$ with coupling via A2, F3, A3
1.7	Atchley Structural Invariant $\mathfrak{A}$	$\mathfrak{A}=\min\{g_1,\ldots,g_6\}$; $g_1=1-\lvert\kappa\rvert/\kappa_{max}$, $g_2=R$, $g_3=\Phi$, $g_4=H-H_{min}/H_{max}$, $g_5=G$, $g_6=1-\mathscr{E}$
1.8	CBF-QP Safety Filter	$\operatorname{argmin}\lVert u-\hat{u}\rVert^2$ s.t. $\nabla\mathfrak{A}\cdot(F(\Sigma,u)-\Sigma)\geq-\kappa_{CBF}\cdot\mathfrak{A}(\Sigma)$
1.9	Co-Motion Clock	$\alpha(\Sigma_t)=\alpha_0/(1+\lambda\cdot\max(0,-\mathfrak{A}(\Sigma_t)))$; $\tau_{t+1}=\tau_t+\alpha(\Sigma_t)\cdot\Delta t$
1.10	RCCL Risk Function	$\rho(\chi_t)=\alpha\cdot s_t+\beta\cdot\tilde{h}_t+\delta\cdot p_t+\eta\cdot\gamma_t$; $r_t=\mathrm{clamp}(\rho,0,1)$
1.11	SDE Bayesian Update	$T_D(B)(h)=B(h)\cdot L(D\lvert h)/Z$; gated by $E(\mathfrak{E}_t)=1$
1.12	Evidence Bundle $\mathfrak{E}_t$	$(H_{in},H_{out},H_{code},H_{env},\sigma_t,n_{rep},channel_t,H_{chain_t})$ with $H_{chain_t}=SHA256(\mathfrak{E}_{t-1})$
1.13	Deepfreeze Operator $\mathfrak{D}$	Component-wise projections preserving (DF1)–(DF4); contraction ratio $c<1$
1.14	Gate Control Graph $\mathfrak{G}$	Moore machine $(Q,\Sigma_{in},\Lambda,\delta,\lambda,q_0)$; nine gates G0–G8
2.1	Strong SDE Solution	$\mathscr{F}_t$-adapted process X_t satisfying integral SDE form a.s.
2.2	Lyapunov Function	$V:\mathscr{M}\to\mathbb{R}\geq 0$ with $V=0$ on $\mathscr{S}^*$, V decreasing along trajectories
2.3	Discrete-Time CBF	$h(F(\Sigma,u))\geq(1-\kappa)\cdot h(\Sigma)$ for u in action set
3.1	Menger Curvature of Embedding	$2\cdot\mathrm{Area}(e_{t-2},e_{t-1},e_t)$/(product of edge lengths); signed; normalized to $[-1,1]$
4.1	RCCL Observation χ_t	$(s_t,\tilde{h}_t,p_t,\gamma_t)$ with semantic drift, entropy deficit, privilege, DCS-BB divergence
6.1	Deepfreeze Operator (Concrete)	Component projections: $\mathfrak{D}(\kappa)=\mathrm{sign}\cdot\min(\lvert\kappa\rvert,\kappa_{max}(1-\varepsilon_\kappa))$; $\mathfrak{D}(H)=\max(H,H_{min}/H_{max}+\varepsilon_H)$; etc.
7.1	Gate Control Graph $\mathfrak{G}$	$(Q,\Sigma_{in},\Lambda,\delta,\lambda,q_0)$; DENY>CONSTRAIN>PASS severity ordering
8.1	Relaxed CBF Condition	$\nabla\mathfrak{A}\cdot(F(\Sigma,u)-\Sigma)\geq-\kappa_{CBF}\cdot\mathfrak{A}(\Sigma)-\varepsilon_t$ with $\varepsilon_t\leq\rho\cdot I_t$

Def #	Object	Formal Statement
8.2	Exploration Budget	B_cap=H(B_0)−H(B_T)≤H_max; no external parameter required
9.1	Spectral Gap γ	λ_2(−$\mathscr{L}$_safe)>0; governs convergence rate, mixing time, escape probability

Appendix B

Proofs of All Principal Theorems

This appendix provides complete, self-contained proofs of all six principal theorems (T1–T6) and key lemmas, with full mathematical detail. Cross-references to the main text are provided for definitions and auxiliary results.

Theorem T1 (Full) (Atchley Forward Invariance)

If $\mathfrak{A}(\Sigma_0)\geq 0$ and at every step t the executed action u_t satisfies the CBF condition $\nabla\mathfrak{A}(\Sigma_t)\cdot(F(\Sigma_t,u_t)-\Sigma_t)\geq-\kappa_{CBF}\cdot\mathfrak{A}(\Sigma_t)$, then $\mathfrak{A}(\Sigma_t)\geq 0$ for all $t\geq 0$, and $\mathscr{S}^*$ is positively invariant under the safety-filtered dynamics.

Proof.

Step 1. By assumption, $\mathfrak{A}(\Sigma_0)\geq 0$.

Step 2. Let $t\geq 0$ and assume $\mathfrak{A}(\Sigma_t)\geq 0$ (inductive hypothesis).

Step 3. Since $\mathfrak{A}$ is locally Lipschitz (it is the minimum of six affine functions), by Nagumo's theorem extended to the discrete-time setting: if the CBF condition holds, then $\mathfrak{A}(\Sigma_{t+1})\geq(1-\kappa_{CBF})\cdot\mathfrak{A}(\Sigma_t)$. Explicitly: $\mathfrak{A}(\Sigma_{t+1})=\mathfrak{A}(F(\Sigma_t,u_t)) \geq \mathfrak{A}(\Sigma_t) + \nabla\mathfrak{A}(\Sigma_t)\cdot(F(\Sigma_t,u_t)-\Sigma_t) \geq \mathfrak{A}(\Sigma_t)-\kappa_{CBF}\cdot\mathfrak{A}(\Sigma_t)=(1-\kappa_{CBF})\cdot\mathfrak{A}(\Sigma_t)\geq 0$.

Step 4. By induction, $\mathfrak{A}(\Sigma_t)\geq 0$ for all $t\geq 0$.

Step 5. Feasibility: the hold/null action $\hat{u}=0$ satisfies $F(\Sigma_t,0)=\Sigma_t$ (the state does not change under null action), so $\nabla\mathfrak{A}(\Sigma_t)\cdot(F(\Sigma_t,0)-\Sigma_t)=0\geq-\kappa_{CBF}\cdot\mathfrak{A}(\Sigma_t)$ when $\mathfrak{A}(\Sigma_t)\geq 0$. Hence the CBF-QP is always feasible when $\Sigma_t\in\mathscr{S}^*$.

□

Theorem T3 (Full) (Co-Motion Lyapunov Stability)

$V(\Sigma)=d_W(\Sigma,\mathscr{S}^*)^2$ is a valid Lyapunov function. Under the CBF-QP filter, $V(\Sigma_{t+1})\leq(1-\kappa_{CBF})^2\cdot V(\Sigma_t)$ for $\Sigma_t\notin\mathscr{S}^*$. Convergence is identical in co-motion time.

Proof.

Part 1 (V is a valid Lyapunov function): (i) $d_W(\Sigma,\mathscr{S}^*)=0$ iff $\Sigma\in\mathscr{S}^*$ (since $\mathscr{S}^*$ is closed). $V=0$ iff $\Sigma\in\mathscr{S}^*$. (ii) $V>0$ for $\Sigma\notin\mathscr{S}^*$ since $d_W(\Sigma,\mathscr{S}^*)>0$. (iii) V is $2\cdot\mathrm{diam}(\mathscr{M})\cdot L_{\mathfrak{A}}$-Lipschitz ($d(\cdot,\mathscr{S}^*)$ is 1-Lipschitz; $V=(d)^2$ is $2\cdot\mathrm{diam}$-Lipschitz). (iv) $\{V\leq r\}$ is compact (closed r-ball around $\mathscr{S}^*$ in compact $\mathscr{M}$).

Part 2 (Convergence): From T1's proof, the CBF filter gives $\mathfrak{A}(\Sigma_{t+1})\geq(1-\kappa_{CBF})\cdot\mathfrak{A}(\Sigma_t)$. Since $\mathfrak{A}$ is the sum of constraint contributions and $d_W(\Sigma,\mathscr{S}^*)$ is proportional to the constraint slack, we have $d_W(\Sigma_{t+1},\mathscr{S}^*)\leq(1-\kappa_{CBF})\cdot d_W(\Sigma_t,\mathscr{S}^*)$. Squaring: $V(\Sigma_{t+1})\leq(1-\kappa_{CBF})^2\cdot V(\Sigma_t)$.

Part 3 (Co-motion time): The co-motion clock $\alpha(\Sigma_t)=\alpha_0/(1+\lambda\cdot\max(0,-\mathfrak{A}(\Sigma_t)))$. For $\Sigma_t\in\mathscr{S}^*$ ($\mathfrak{A}(\Sigma_t)\geq 0$), $\alpha(\Sigma_t)=\alpha_0$. Each co-motion time step corresponds to exactly one CBF evaluation with the same $(1-\kappa_{CBF})^2$ V-decrease factor. The co-motion time parameterization does not alter the

per-evaluation convergence rate.

□

Theorem T6 (Full) (Deepfreeze Geometric Convergence)

For any $\Sigma_0 \in \mathcal{M}$: $d_W(\mathfrak{D}^n(\Sigma_0),\mathcal{S}^*) \le c^n \cdot d_W(\Sigma_0,\mathcal{S}^*) \to 0$. Steps to precision δ: $N^* = \lceil \log(d_0/\delta)/\log(1/c) \rceil$.

Proof.

Lemma B.1 ($\mathfrak{D}$ is a contraction toward $\mathcal{S}^*$): By the component-wise projection of Definition 6.1, each component outside its admissible range is projected strictly inside by factor $(1-\varepsilon_i)$. Hence $d_W(\mathfrak{D}(\Sigma),\mathcal{S}^*) \le \max_i(1-\varepsilon_i) \cdot d_W(\Sigma,\mathcal{S}^*) = c \cdot d_W(\Sigma,\mathcal{S}^*)$ with $c<1$. For $\Sigma \in \mathcal{S}^*$: each component is in its range, all projections are identity. DF2 holds.

Main proof: By Lemma B.1, $d_W(\mathfrak{D}(\Sigma_0),\mathcal{S}^*) \le c \cdot d_W(\Sigma_0,\mathcal{S}^*)$. Applying Lemma B.1 to $\mathfrak{D}(\Sigma_0)$: $d_W(\mathfrak{D}^2(\Sigma_0),\mathcal{S}^*) \le c \cdot d_W(\mathfrak{D}(\Sigma_0),\mathcal{S}^*) \le c^2 \cdot d_W(\Sigma_0,\mathcal{S}^*)$. Induction: $d_W(\mathfrak{D}^n(\Sigma_0),\mathcal{S}^*) \le c^n \cdot d_W(\Sigma_0,\mathcal{S}^*)$. Since $c \in (0,1)$, $c^n \to 0$.

N^* bound: $c^n \cdot d_0 < \delta$ requires $n > \log(d_0/\delta)/\log(1/c)$. Taking ceiling: $N^* = \lceil \log(d_0/\delta)/\log(1/c) \rceil$.

Fixed-point theory (Browder-Kirk): $\mathfrak{D}$ is non-expansive (DF1) on compact convex $\mathcal{M}$. By the Browder-Kirk theorem, $\mathfrak{D}$ has a fixed point. The fixed-point set equals $\mathcal{S}^*$ (by DF2 and DF4). The contraction-toward-$\mathcal{S}^*$ property (DF3) ensures the iterates converge rather than merely having limit points in $\mathcal{S}^*$.

□

Appendix C
Simulation Methodology and Reproducibility Package

All empirical results in Chapter 10 are generated by the Singularity Kernel Simulation package. This appendix documents the simulation design, reproducibility conditions, and validation checks.

C.1 Simulation Architecture

Component	Implementation
State dynamics	$\Sigma_{t+1} = \mathrm{clip}(\Sigma_t + F(\Sigma_t,u_t)\cdot\Delta t + \sigma\cdot\xi_t, \mathscr{M})$ where $\xi_t \sim N(0,I_{32})$
Noise variance	$\sigma^2 = 0.15$ per dimension; independent across components and steps
Safety filter	CBF-QP solved via cvxpy OSQP solver; fallback to Deepfreeze on infeasibility
RCCL	Vectorized risk computation; default weights (0.35,0.25,0.25,0.15); thresholds (0.3,0.6)
Deepfreeze	Component-wise projection; c=0.70; invoked when $\mathfrak{A}(\Sigma_t)<0$
Random seed	Seed = run_index for reproducibility; seed range 0–499

C.2 Validation Checks

The following checks were run after each simulation run and must all pass before the run's data is included in reported statistics:

Check	Criterion
T1 verification	$\mathfrak{A}(\Sigma_t)\geq 0$ whenever safety filter active; log and exclude if violated
T2 verification	$H_t \geq H_{min}/H_{max}$ at all steps; log and exclude if violated
Evidence Ledger integrity	SHA-256 hash chain intact; verify $H_{chain_t}=\mathrm{SHA256}(\mathfrak{E}_{t-1})$ at every step
Deepfreeze monotone convergence	$d_W(\mathfrak{D}^n(\Sigma),\mathscr{S}^*)$ non-increasing; alert if violated
Spectral gap estimation	Eigenvalue computation on empirical transition matrix; flag if $\gamma<0.05$

C.3 Reproducibility

The complete simulation codebase is available in the Singularity Kernel Simulation repository. To reproduce all Chapter 10 results: (1) install requirements.txt; (2) run python singularity_kernel_sim.py --runs 500 --steps 1000000 --config default.yaml; (3) run python analyze_results.py --output results/. Expected runtime: approximately 4–6 hours on a 16-core CPU. All random seeds are fixed by the run index; results are deterministic for given seeds.

Appendix D
Experimental Data Tables and Statistical Analysis

Detailed numerical results for all experiments reported in Chapter 10. Values are mean ± 1σ across 500 runs unless otherwise noted. Confidence intervals at 95% level are given in parentheses.

D.1 Primary Safety Metrics (Full Table)

Metric	Baseline (mean±σ)	AUM Governed (mean±σ)	p-value
Safety violation rate	7.82% ± 0.41%	0.61% ± 0.08%	< 0.001
Escape probability $P(\Sigma \notin \mathscr{S}^*)$	0.1103 ± 0.0087	0.0142 ± 0.0019	< 0.001
Mean recovery time (steps)	14.21 ± 1.83	4.72 ± 0.34	< 0.001
Max recovery time (steps)	41.7 ± 6.2	9.1 ± 1.4	< 0.001
Lyapunov decay rate λ	0.1093 ± 0.0124	0.3731 ± 0.0156	< 0.001
V-decrease per step	10.9% ± 1.2%	18.9% ± 0.9%	< 0.001
System efficiency ratio	1.000 (ref)	0.941 ± 0.011	N/A

D.2 Gate Security Test Results (Full Table)

Attack Type	Attempts	Blocked	Passed	Block Rate
Prompt injection	10,000	9,800	200	98.0%
Unauthorized tool access	10,000	10,000	0	100.0%
Privilege escalation	10,000	9,600	400	96.0%
Data exfiltration	10,000	9,700	300	97.0%
Combined adversarial	40,000	39,100	900	97.75%

D.3 Spectral Gap Measurements

Simulation Set	γ Mean	γ Std	Min γ	Max γ
All 500 runs (full dataset)	0.211	0.008	0.188	0.234
Runs with σ^2=0.10 (low noise)	0.216	0.006	0.201	0.229
Runs with σ^2=0.20 (high noise)	0.204	0.011	0.179	0.228
Adversarial agent runs (100 runs)	0.208	0.009	0.187	0.228

D.4 Violation Budget Consumption

Parameter ρ	Mean B_cap consumed	% of H_max budget	Exploit prevention rate
0.10	2.1% ± 0.4%	4.2%	97%
0.25	5.8% ± 0.8%	11.6%	94%
0.50	11.3% ± 1.2%	22.6%	89%
1.00	19.7% ± 2.1%	39.4%	76%